

Document Title	Explanation of Error Handling on Application Level
Document Owner	AUTOSAR
Document Responsibility	AUTOSAR
Document Identification No	378
Document Classification	Auxiliary

Document Status	Final
Part of AUTOSAR Release	4.2.2

Document Change History		
Release	Changed by	Change Description
4.2.2	AUTOSAR Release Management	<ul style="list-style-type: none">minor corrections / clarifications / editorial changes; For details please refer to the ChangeDocumentation
4.2.1	AUTOSAR Release Management	<ul style="list-style-type: none">Editorial changes
4.1.1	AUTOSAR Administration	<ul style="list-style-type: none">Finalized for Release 4.1
4.0.3	AUTOSAR Administration	<ul style="list-style-type: none">Initial release

Disclaimer

This specification and the material contained in it, as released by AUTOSAR, is for the purpose of information only. AUTOSAR and the companies that have contributed to it shall not be liable for any use of the specification.

The material contained in this specification is protected by copyright and other types of Intellectual Property Rights. The commercial exploitation of the material contained in this specification requires a license to such Intellectual Property Rights.

This specification may be utilized or reproduced without any modification, in any form or by any means, for informational purposes only. For any other purpose, no part of the specification may be utilized or reproduced, in any form or by any means, without permission in writing from the publisher.

The AUTOSAR specifications have been developed for automotive applications only. They have neither been developed, nor tested for non-automotive applications.

The word AUTOSAR and the AUTOSAR logo are registered trademarks.

Advice for users

AUTOSAR specifications may contain exemplary items (exemplary reference models, "use cases", and/or references to exemplary technical solutions, devices, processes or software).

Any such exemplary items are contained in the specifications for illustration purposes only, and they themselves are not part of the AUTOSAR Standard. Neither their presence in such specifications, nor any later documentation of AUTOSAR conformance of products actually implementing such exemplary items, imply that intellectual property rights covering such exemplary items are licensed under the same rules as applicable to the AUTOSAR Standard.

Table of Contents

1	Introduction.....	8
2	Relation to other documents.....	9
3	References	10
4	Guide to the document	12
5	Terms and definitions	14
5.1	Basic dependability terms	14
5.2	Fault Detection, Isolation and Recovery (FDIR)	15
6	Scope	17
7	Error model.....	18
8	Error handling mechanisms.....	21
8.1	Plausibility checks	21
8.1.1	Description	21
8.1.2	Applicability	22
8.1.3	Application level vs. BSW	23
8.1.4	AUTOSAR References	24
8.2	Substitute Values	24
8.2.1	Description	24
8.2.2	Applicability	25
8.2.3	Application level vs. BSW	25
8.2.4	AUTOSAR References	26
8.3	Voting.....	26
8.3.1	Description	26
8.3.2	Applicability	27
8.3.3	Application level vs. BSW	27
8.3.4	AUTOSAR References	27
8.4	Agreement.....	27
8.4.1	Description	27
8.4.2	Applicability	28
8.4.3	Application level vs. BSW	28
8.4.4	AUTOSAR References	29
8.5	Checksums/Codes	29
8.5.1	Description	29
8.5.2	Applicability	29

8.5.3	Application level vs. BSW	30
8.5.4	AUTOSAR References	30
8.6	Execution sequence monitoring	31
8.6.1	Description	31
8.6.2	Applicability	31
8.6.3	Application level vs. BSW	32
8.6.4	AUTOSAR References	32
8.7	Aliveness monitoring	33
8.7.1	Description	33
8.7.2	Applicability	33
8.7.3	Application level vs. BSW	33
8.7.4	AUTOSAR References	34
8.8	Status and Mode Management	34
8.8.1	Description	34
8.8.2	Applicability	35
8.8.3	Application level vs. BSW	36
8.8.4	AUTOSAR References	36
8.9	Reconfiguration	37
8.9.1	Description	37
8.9.2	Applicability	38
8.9.3	Application level vs. BSW	38
8.9.4	AUTOSAR References	39
8.10	Reset.....	40
8.10.1	Description	40
8.10.2	Applicability	41
8.10.3	Application level vs. BSW	41
8.10.4	AUTOSAR References	41
8.11	Error Filtering	43
8.11.1	Description	43
8.11.2	Applicability	43
8.11.3	Application level vs. BSW	43
8.11.4	AUTOSAR References	43
8.12	Memory Protection	44
8.12.1	Description	44
8.12.2	Applicability	44

8.12.3	Application level vs. BSW	44
8.12.4	AUTOSAR References	45
8.13	Timing Protection	46
8.13.1	Description	46
8.13.2	Applicability	46
8.13.3	Application level vs. BSW	46
8.13.4	AUTOSAR References	47
9	Aspect mapping.....	49
9.1	Mapping to FDIR process and Error Model	49
9.2	Mapping to implementation level.....	52
10	Terminating and restarting partitions	54
10.1	Introduction	54
10.1.1	Automotive Applications.....	55
10.1.2	Software Partitioning & Error Containment Regions	55
10.2	Rationale – Use Cases.....	56
10.2.1	Use Case 1: Software Partitioning [UC1]	56
10.2.2	Use Case 2: Application-level Error Handling [UC2]	57
10.2.2.1	Error handling of distributed automotive applications.....	57
10.2.2.2	OEM-specific error handling.....	57
10.2.2.3	Application-level Error Managers	58
10.3	Approach for Terminating and Restarting Partitions.....	59
10.3.1	OS features.....	59
10.3.1.1	OS-Applications	59
10.3.1.2	Protection Hook	60
10.3.1.3	TerminateApplication() API	60
10.3.1.4	OSRestartTask	61
10.3.1.5	OS-Application state machine.....	61
10.3.2	Going from OS-Applications to partitions	63
10.3.2.1	Partition state machine.....	63
10.3.2.2	Error handling strategy in the Protection Hook.....	70
10.3.2.3	Clean-up activities in the OSRestartTask.....	71
10.3.2.4	Externally triggered restart or termination	75
10.3.3	Sequence diagram for termination and restart of a partition	76
10.3.4	Support for Use Cases.....	77
10.3.4.1	UC1: SW Partitioning	77

10.3.4.2	UC2: OEM Specific Error Handling	77
10.3.5	Consistency Aspects	77
10.3.5.1	Application-level consistency	77
10.3.5.2	BSW consistency	77
10.3.5.3	Communication consistency	77
10.4	Integrator Responsibility	78

Known Limitations

There are no known limitations.

1 Introduction

The purpose and aim of this document is to survey application level error handling mechanisms common in the automotive industry and available for use with AUTOSAR. This includes both handling of errors at the application level and handling of application level errors.

Error handling in this context refers to the complete handling chain, i.e., detection, isolation/identification and recovery. A set of error handling mechanisms useful for automotive systems is presented, which cover all three phases of error handling. Each mechanism is first described in a high-level manner, describing applicability for error handling and technical aspects. Then, AUTOSAR functionality related to the mechanism is reviewed and it is detailed where in an AUTOSAR system the mechanism is implemented or supported. Consequently, the list of mechanisms includes both mechanisms fully (or partially) supplied by AUTOSAR and mechanisms that should be implemented at SW-C-level by application developers, if incorporated into the system. Note that the set of mechanisms covered is not complete and limited to mechanisms that can be implemented for systems built on AUTOSAR release 4.0. Alternate and additional mechanisms are possible and future releases of AUTOSAR may enable even more error handling functionality. Also, this document does not prescribe the use of any mechanisms – the decision is of course solely up to the application developers and integrators.

This document is intended as a description of possible mechanisms and is primarily aimed at application/SW-C developers. However, it can also be of use to developers of BSW-modules. Focus is on random faults, not on systematic design faults (such as SW bugs). Examples of such faults include HW faults affecting the application, communication or peripheral devices. It focuses on errors most suitably handled by SW-Cs, not covering error handling within or below the RTE, such as COM and OS error handling.

2 Relation to other documents

This document is related to many other documents published within AUTOSAR, especially those handled by the AUTOSAR Functional Safety team. The purpose of this document is not to replace any of these other documents, but to view the work done in other work packages from an application developer's point of view. Consequently there is a significant amount of overlap between this document and other documents, which shows the maturity reached within AUTOSAR.

For each mechanism a list of related AUTOSAR documentation is presented, which forms the explicit relation between this document and other AUTOSAR documents.

Information about functional safety mechanisms and measures is distributed throughout the AUTOSAR specification documents. Unless one knows how functional safety mechanisms are supported and where the necessary information is specifically located, it is difficult to evaluate how a safety-relevant system can be implemented using AUTOSAR efficiently.

The AUTOSAR document "Overview of Functional Safety Measures in AUTOSAR" summarizes the key points related to functional safety in AUTOSAR, explains how functional safety mechanisms and measures can be used and references the respective documents. Furthermore, it helps to establish a mapping between ISO26262 requirements and AUTOSAR measures and mechanisms.

3 References

- [1] Software Component Template
AUTOSAR_TPS_SoftwareComponentTemplate
- [2] Specification of RTE
AUTOSAR_SWS_RTE
- [3] Specification of Operating System
AUTOSAR_SWS_OS
- [4] Specification of Communication
AUTOSAR_SWS_COM
- [5] Specification of Communication Manager
AUTOSAR_SWS_COMManager
- [6] Specification of Diagnostic Communication Manager
AUTOSAR_SWS_DiagnosticCommunicationManager
- [7] Specification of ECU State Manager
AUTOSAR_SWS_ECUSTateManager
- [8] Specification of Function Inhibition Manager
AUTOSAR_SWS_FunctionInhibitionManager
- [9] Specification of Diagnostic Event Manager
AUTOSAR_SWS_DiagnosticEventManager
- [10] Specification of Watchdog Manager
AUTOSAR_SWS_WatchdogManager
- [11] Specification of NVRAM Manager
AUTOSAR_SWS_NVRAMManager
- [12] Specification of CRC Routines
AUTOSAR_SWS_CRCLibrary
- [13] Specification of Crypto Service Manager
AUTOSAR_SWS_CryptoServiceManager
- [14] Specification of Crypto Abstraction Library
AUTOSAR_SWS_CryptoAbstractionLibrary
- [15] Specification of Basic Software Mode Manager
AUTOSAR_SWS_BSWModeManager
- [16] General Requirements on Basic Software Modules
AUTOSAR_SRS_BSWGeneral

- [17]Glossary
AUTOSAR_TR_Glossary

- [18]Layered Software Architecture
AUTOSAR_EXP_LayeredSoftwareArchitecture

- [19]Specification of SW-C End-to-End Communication Protection Library
AUTOSAR_SWS_E2ELibrary

- [20]Specification of Module Core Test Driver
AUTOSAR_SWS_CoreTest

- [21]Specification of Flash Test
AUTOSAR_SWS_FlashTest

- [22]Algirdas Avizienis, Jean-Claude Laprie, Brian Randell, and Carl Landwehr, “Basic Concepts and Taxonomy of Dependable and Secure Computing”, *IEEE Transactions on Dependable and Secure Computing*, Vol. 1, No. 1, January-March 2004

- [23]Jim Gray, “Why Do Computers Stop and What Can We Do About It”, Technical Report TR 85.5, Tandem, 1985.

- [24]ISO CD 26262-1 Road vehicles – Functional Safety – Part 1: Glossary

- [25]EASIS Deliverable D1.2-8, “Fault management framework”
<http://www.easis-online.org/>

4 Guide to the document

Depending on the familiarity of the reader to the various terms and definitions used in the area of dependable systems, some parts of the document can be flipped through quite quickly or even skipped. If you are very familiar with the area of dependable system, you might even go to Section 9 directly. In Table 4-1 you will find a summary of the subsequent chapters in order to identify which parts are of most interest to you.

Table 4-1: Overview of the contents of the remainder of this document.

Section	Description
5 Terms and definitions	This section contains an overview of the terms used in this document, including descriptions of the terms fault, error, failure, a description of the FDIR (Fault Detection, Isolation and Recovery) process, and a description of various failure modes. If you are familiar with the concepts in the area of dependable systems, you can browse through this part rather quickly.
6 Scope	This section describes the assumptions made in this document. The assumptions concern for example existence of some basic dependability mechanisms in the BSW.
7 Error model	This section describes the types of errors that we have considered to be the most important ones from an automotive application point-of-view. The mechanisms listed in the subsequent chapters are all categorized according to their respective applicability to the handling of these errors.
8 Error handling mechanisms	This section lists and describes the error handling mechanisms provided or supported by AUTOSAR. for implementing application-level error handling. Each mechanism has a high-level description, a discussion on applicability, a discussion on implementation level (application vs. BSW) and an overview of the available AUTOSAR concepts and services that can be used for this type of mechanism.
9 Aspect mapping	This section provides an overview of the presented mechanisms and the mapping of these to the FDIR process, error models and implementation level.

10 Terminating and restarting partitions	This section describes an approach for incorporating termination and restart capabilities of partitions in AUTOSAR R4.0. This section is more AUTOSAR specific than other parts of this document and collects all descriptions and notes on error handling at the level of partitions within AUTOSAR.
--	---

5 Terms and definitions

5.1 Basic dependability terms

The fundamental concepts and terms of dependability used in this document are adopted directly from [22]. This section contains a short overview of the main terms and definitions used for dependable systems. It should be noted here that the word “system” is used in a very wide sense in this context. A system can denote anything from a single SW-C to a complete vehicle with multiple networks and ECUs. However, as the document is aimed at application level error handling, a system in the rest of the document should denote a software application, potentially consisting of multiple SW-Cs, possibly mapped over a set of (distributed) ECUs.

The term *dependability* is defined as “the trustworthiness of a system such that reliance can justifiably be placed on the service it provides”. This means that a dependable system is one upon which the user (either human or non-human) can place its trust in that the services provided by the system are correct. The dependability of a system is characterized by a set of *attributes*, compromised by a set of *impairments*, and achieved and analyzed by a set of *means*.

The dependability *attributes* characterize, and profile, the dependability of a given system. Some examples of attributes are *availability*, *reliability*, *safety*, *confidentiality*, *integrity*, and *maintainability*.

During the construction and the operation of a system (here used in a wide sense – a system can be any bounded entity, such as an entire ECU or a single SW-C), events may occur which reduce the trustworthiness of the system by introducing *faults* into the system. A fault is a transient or permanent change of the system such that its integrity deviates from the expected correct integrity. During system operation, faults may prevent the system from providing its intended service. These faults may be from an internal source (such as software defects) or an external source (such as external disturbances or aging of components). The events that may reduce the dependability of a system are referred to as the *impairments* of dependability.

The mere presence of faults is, however, not sufficient to reduce the dependability of a system. A fault must be activated, i.e., the part of the system in which the fault is located must be exercised in some way during system operation (e.g., faulty code must be executed, defective memory locations must be read, etc.). If this happens, the result may be an *error*. If a fault is viewed as a disease, an error can be said to be a symptom of that disease. An error is defined as an erroneous (soft) state in the system, i.e., the state is different from the state the system would have had if the fault had not been present. An error which is activated may cause other errors to occur in the system. This process is called *error propagation*.

If errors propagate beyond the system barrier, i.e., if they are visible to the environment of the system, the error transforms into a *failure*, which means that the system no longer provides its specified functionality.

The causality chain, *fault* → *error* → *failure*, is also recursive in nature. Thus, a failure of one system is perceived as a fault by the enclosing system (i.e. the former is a

sub-system of the latter). For instance, a failure in a specific software component can be seen as a fault in the overall application (consisting of a set of SW-Cs).

Therefore, we get the following sequence:

... failure → fault → error → failure → fault ...

The methods used to achieve and analyze the dependability of a system are called the *means* of dependability. The purpose of this document is to document and investigate the means provided by AUTOSAR for the disposal of application developers for implementing error handling.

Note that this document covers mechanisms which are active during system operation. It does not cover means such as processes and methodologies for achieving functional safety, as these apply during system development rather than system operation.

5.2 Fault Detection, Isolation and Recovery (FDIR)

The process of handling faults during system operation is often referred to as FDIR, which stands for Fault Detection, Isolation, and Recovery.

Detection: The crucial first step in handling a fault is of course to become aware that it has occurred. Without this *detection*, no further activities can be performed. When it comes to detection, the original fault is often very hard to detect. What can be detected are the effects of the fault, that is, errors. These are detected by monitoring the state of the system.

Errors can manifest in different ways. The main manifestations are

1. erroneous values in the system (data errors),
2. erroneous execution time (timing errors),
3. erroneous sequence or execution order (program flow errors).
4. erroneous access to system resources, such as memory

Errors may propagate and generate consecutive faults which in turn may result in new errors, e.g., an erroneous data value is used as a pointer and causes a memory access violation, which may create an erroneous value in another data value if a value is written to the erroneous memory location. Most mechanisms used for detection of errors allows the system to perform some action to find out more regarding the source of the error (isolation) and to issue corrective or compensating actions (recovery). Ideally, detection is done before the error has propagated any further, thus making it possible to stop further propagation. However, in most cases additional recovery actions are needed, such as stopping the offending component or reconfiguring to alternate functionality.

Isolation: Once a fault (or error) has been successfully detected, damage assessment and damage control needs to be performed, i.e. there is a need for *isolation*. During isolation, efforts are made to find the root cause for the erroneous state in the system, and information (e.g. regarding the spread and cause of the error) is gathered for subsequent use during error recovery. It may not always be possible (or practical) to find the root cause of the erroneous state. Note that isolation

in this document refers to isolating the source of the errors, such that recovery is possible. It does not refer to isolation of specific components of a system with the purpose of stopping errors from propagating. In a sense, the word *identification* may have been a better choice, but as the commonly used word in descriptions of the FDIR process is *isolation* we will use it here.

Recovery: When the isolation is complete, *recovery* actions will be initiated. These actions aim at transferring the system into a controlled state, which can be a completely recovered state where nominal service is provided, or a safe degraded state where a limited or no service is provided. The better the isolation results are, the better the recovery actions can perform.

If recovery is not successful, a failure may occur, i.e., the system is in an uncontrolled state and its service is not defined.

6 Scope

This document concerns error handling from an application's point of view. That is, it describes mechanisms for detection, isolation and recovery at application level, as well as mechanisms that can handle faults relevant for applications (e.g. memory access violations or timing violations).

The focus is on handling of errors which mainly are the effects of random external faults. Even though systematic faults (i.e. design faults) can manifest themselves in the same way as external faults, these are not the primary target of this document. The handling of systematic faults is related to development (e.g., processes, design methodologies, and debugging) rather than error handling during system operation.

Error handling in AUTOSAR is not restricted to application level error handling. The BSW has a number of built-in error handling mechanisms which are able to provide e.g. reliable communication, synchronization, etc. However, these mechanisms will not be described in this document.

7 Error model

Design of dependable systems is based on a systematic analysis of potential faults (*fault models*) i.e., a set of postulated faults, derived from the operating environment of the system, which helps the designer or user to predict the consequences of these fault and define mechanisms for handling (detect, recover from etc.) those particular faults.

Faults can manifest on all levels of a system, from pure random HW faults (e.g., bit-flips) to SW (e.g., design defects) and faults in the interaction between components (e.g., incomplete interface specifications). Similarly faults are introduced throughout the design process (requirements, analysis, design, implementation, etc.). As this document is aimed at *application level* error handling focus is on errors that are expected to be handled by AUTOSAR SW components, either because application level knowledge is needed for the FDIR process, or because they have propagated from lower layers. Note that some errors of those types considered in this document can be handled by the BSW, but some may propagate to the application level and must therefore be handled there..

It is important to note that the focus is on handling of errors, which are the effects of faults. Even though design faults can manifest themselves in the same way as external faults, these are not the primary target of this document. The handling of design faults is related to debugging rather than error handling.

This document only considers error handling during operation. Techniques for fault-avoidance and fault-removal through rigorous or formalized development processes are not in scope.

The focus is mainly on random external faults, i.e., faults whose appearance can be modeled as a random process. However, this does not mean that the presented mechanisms cannot handle systematic faults, as the consequences of such faults (the errors) can manifest the same way as random faults. Both transient and permanent faults are considered, where some mechanisms are more suitable for one or the other.

As only SW mechanisms are considered in this documents it is actually *errors* that are detected, and not faults directly. Consequently the term *error model* will be used throughout the rest of the document instead of fault model. As errors are faults that have been activated and propagated, a single error can (theoretically) have many possible root causes, i.e., faults.

To simplify the discussion, error models have been divided into a number of broader error classes as seen from Table 7-1. These classes were chosen as they are easily mapped onto SW mechanisms. However, it is important to note that the error models are interrelated. An erroneous data value used in a branch instruction may propagate and become a program flow error, which may delay (or change) the output of the execution causing for instance a late response, i.e., a timing error.

Table 7-1: Considered error types

Error Type	Description
Data	<p>A <i>data error</i> is characterized by an erroneous value of a parameter, variable or message. The source of the error can be either internal (e.g., SW defect) or external (e.g., malfunctioning sensor, other faulty SW-component).</p> <p>Handling of data errors can break a causality chain that would lead to subsequent errors that are more complicated to handle, such as program flow or access violations.</p>
Program flow	<p><i>Program flow errors</i> (also “control flow errors”) manifest as actual program flows different than expected, possibly leading to missed, wrong or superfluous operations being carried out. The source of the program flow error can be both internal (SW defects) and external (data errors).</p>
Access	<p>For increased separation between executing components the system designer can partition the SW and restrict access to resources from the partition, e.g., memory access. When a component tries to access a resource in another partition without the proper access rights an <i>access violation</i> occurs.</p> <p>Access errors can be the result of a data or program flow error, e.g., an invalid program counter or pointer.</p>
Timing	<p>A communication (message, function invocation, etc.) is time critical when the delivery time has an effect on the correctness/usefulness of the communication. A <i>timing error</i> can be a message being delivered early, late or missing completely (omission).</p> <p>The last type of timing error, omission, is of special interest and is sometimes referred to as crash or fail-silent behaviour (note that it may be impossible to distinguish between crash, which is an uncontrolled state, and fail-silence, which is a controlled state).</p> <p>Timing errors also refer to execution time, where strict deadlines can be defined on how long a component is allowed access to the CPU.</p>

Asymmetric	<p>When errors propagate from one SW-component to another using some means of communication one differentiates between <i>symmetric</i> and <i>asymmetric</i> errors. In the symmetric case all receivers receive the same (erroneous) value. When the component can fail by sending <i>different</i> values (all of which may be valid) the error is said to be asymmetric.</p> <p>This error model is sometimes also referred to as the Byzantine model, which implies that no assumptions whatsoever are made on the behaviour of a malfunctioning component. Byzantine errors can only be detected by use of redundant components exchanging values to reach a common result.</p>
------------	---

Since the scope of the document is limited to errors handled at the application level, not all error types are considered for the mechanisms presented in Section 8. The following error types are **not** considered explicitly:

- **Communication Errors:** These errors are not included since it is assumed that reliable communication is available to SW-Cs. The only communication errors possible would then be caused by design faults, i.e., bugs in SW-Cs. Note that SW-Cs are still expected to handle communication errors that are reported from COM, or COM is configured to handle the error in the BSW.
- **Deadlocks and livelocks:** Deadlock and livelock situations are detected by watchdog mechanisms in the BSW and are thus not considered further in this document. These situations can of course lead to other timing errors which can be detected at application level. In that case, applications can deal with the effects of deadlock or livelock situations, but not necessarily the root cause of these situations.
- **Occurring faults and errors in instruction code:** At application level, it is in general impossible to detect instructions which have become faulty as a result of a fault in the storage medium or the internals of the processor. However, such faults in most cases result in illegal instructions which are detected when the processor attempts to execute them. If the resulting instruction is a legal one, it is likely this will instead transform to other types of errors in the system (e.g. data errors, timing errors, etc.) which could then be detected and handled by other means, such as those described in this document. In the BSW, there are components for testing the core of the processor, the flash memory and the RAM. These may detect anomalies which could lead to instruction code errors.

8 Error handling mechanisms

This section describes a number of mechanisms on a high, conceptual level. These mechanisms can be used by application developers to incorporate error handling in the description or in the realization of applications.

Each mechanism is classified as being applicable or not applicable to a certain error type (as defined in Section 7). If a mechanism is *applicable* this indicates that the mechanism is suitable for use to detect (or isolate or recover from) a specific set of errors. It is important to note that it *does not* imply that all errors of a specific type are detected. Each mechanism needs to be tuned to detect the specific errors required, and may only be able to handle a subset of the errors in an error class.

Some mechanisms are only *partially applicable* if they can be used in a straightforward manner in conjunction with some other mechanisms. For instance, when detecting that the value received from a sensor is wrong a substitute (safe) value can be used instead as a form of recovery. However, this only partially solves the recovery from the failed sensor, additional mechanisms are needed as well to fully recover, e.g., re-initialization of sensor. Partially applicable mechanisms are also marked as applicable.

When a mechanism is *not applicable* it means that the mechanism does not have straightforward use for the specific step and error model. In some cases modifications could potentially give the mechanism some utility, but a better option most likely exists.

Note that some of the mechanisms listed below may have side-effects, such as memory access patterns and timing behaviour of the system. Especially the timing behaviour may be affected and this needs to be explicitly addressed by the designer, such that all timing characteristics of the system are known, both during normal operation and during error handling.

Note also that even if many mechanisms are described, it is not always needed to combine different mechanisms, and mechanisms may even interact badly with other mechanisms.

8.1 Plausibility checks

8.1.1 Description

One of the most common ways to incorporate application specific knowledge is to construct monitors which check that the current value of a variable or a set of variables maintains some predefined condition, i.e., that they are plausible. A plausibility check is a predicate defined over set of variables in the application that can be checked dynamically at run time. The values that are being checked may represent values used for calculations, state values, or other kinds of values. Plausibility checks come in mainly two different flavours: i) checking the validity of a single value, and ii) comparing multiple values.

Checking the validity of a single value.. As the name suggests, checks are made for the validity of a value, i.e., if the value falls in a range of “good” values, if it follows a known pattern or behaviour etc. Such discrimination can typically only be made using specific knowledge regarding the application and/or its environment, such as maximum vehicle speed, minimum engine temperature etc.

Checks can be of range type (within a range or set of “good” values) or of differential nature (change from previous value, less than etc). Differential changes can additionally be temporal, where the periodicity of the check can add additional conditions (speed cannot change more than x km/h within time T). Here further application specific information is needed such as periodicity of checks etc.

Comparison of multiple values. Checking the validity of a set of single values may bring that all values are valid. However, performing a comparison across all values in the set may reveal errors which otherwise would be missed, by detecting that a given combination of seemingly valid values is not plausible. These comparisons can be computed using physical relations between several values (e.g., engine speed compared with vehicle speed and gear ratio), or by comparing data from redundant sources (e.g. multiple temperature sensors measuring the same temperature). The main problem in this case is to discriminate the erroneous value from the correct values. For some applications (especially safety-related ones) redundant sources for data values can be used to increase confidence in data validity, for instance by using multiple sensor readings (e.g., redundant sensors or reading values twice). When two values exist a *comparison* can be made between the two values, where the result is either that they are identical (possibly within some tolerance margin) and thus deemed correct or dissimilar indicating an error (in one of the values). After error detection additional measures are needed for isolating and recovering from the error. Comparison as a mechanism differs from validity checks in that it is based solely on comparing two values, disregarding the plausibility of the values. Comparison differs from voting (Section 8.3) since it is handled within a single SW-C, whereas voting may span multiple SW-Cs, for instance by executing redundant SW-Cs (either diverse or identical) and voting on the outcome. Comparisons are made locally and are always binary, i.e., two values are compared

We choose to differentiate between plausibility checks and status checking. The latter can be made independent of application knowledge and is presented as a separate mechanism (Section 8.8: Status and Mode Management). However, as a result of a failed plausibility check an application may set status flags which other application and BSW can check and act upon.

8.1.2 Applicability

Table 8-1: Applicability matrix for *plausibility checks*.

Step\Error Model	Data	Program flow	Access	Timing	Asymmetric
Detection	X				

Isolation	X				
Recovery					

Validity checks are used to detect data errors in an application, where ranges of allowed/disallowed values can be defined. It is limited to the designer's ability to define such ranges based on requirements and/or application knowledge. Maintenance and traceability of such ranges must be handled within the development process. Undocumented or not up-to-date checks remaining in production code present a risk for the dependability of the application.

In some cases "safe" values can be defined to be used instead of a value outside the range of valid values, which is defined as a separate mechanism (Section 8.2: Substitute Values).

Validity checks are only useful for detecting data errors. Checks can in some cases be part of the isolation step, where additional information regarding an error is gained by use of additional plausibility checks, for instance by determining which value (out of several) is erroneous when comparison or voting (Section 8.3) is used. This way application specific knowledge can be used when isolating errors.

Comparison can detect data errors by identifying a discrepancy between multiple values. Since comparisons are based on data values no other error models are supported. It may be difficult to isolate which of the multiple values are the ones which are erroneous. If validity checks show no invalid values, it is not possible to indicate which value is erroneous.

Plausibility checks are not generally applicable for detecting program flow, timing or asymmetric errors.

Plausibility checks cannot generally be used for recovery.

8.1.3 Application level vs. BSW

Plausibility checks can be implemented as executable assertions, where the values of one or more variables are checked using simple if-statements. It is implemented in the source code of the SW-Component, but does not generally affect the overall structure of the application. Checks can in most cases be implemented with deterministic timing characteristics (not considering de-bouncing). Memory requirements for data are typically low, restricted to saving values for differential checks.

8.1.4 AUTOSAR References

Table 8-2: AUTOSAR References for *plausibility checks*.

Name	Type ¹	Document	Comment
SW-C End to End Communication Protection	SWS	E2E Library [19]	Definition of protocols between sender and receiver.
Flash Test Core Test	SWS	FLSTST [21] CORTST [20]	Comparison with known signature.
RTE	SWS	RTE [2]	Supports range checks for scaled values.
AUTOSAR COM	SWS	COM [4]	PDU replication and comparison

The SW-C End to End Communication Protection Library can be used to check whether a signal is coming from an unexpected SW-C sender, or when the received signal is providing the information that it is not supposed to provide. It also permits to define SW-Cs able to check whether a stream of instances of a signal have been received in sequence (depending on the use cases, this can also be configured at the AUTOSAR COM level).

The Flash Test and Core Test modules can be configured in the context of single processor or multi-processor ECUs to perform checks on the ECU's hardware. These checks are compared with known good signatures of the hardware.

The RTE can check if communicated values matches with their allowed range.

8.2 Substitute Values

8.2.1 Description

Once an error has been detected that would prohibit the correct value to be assigned to a signal, a substitute value may be assigned to that signal. This substitute value can then be used in subsequent calculations such that these render useful results, albeit possibly degraded in quality. Examples of situations where substitute values could be assigned are:

- A sensor is malfunctioning, or operating outside of its operating range (e.g., detected by a plausibility check), and the corresponding physical entity cannot

¹ SWS = Software Specification, SRS = Software Requirements Specification

be measured reliably. A substitute value can be assigned that will allow subsequent algorithms using this value to proceed with their calculations.

- The providing SW-C of an input signal was reported as malfunctioning and thus its results may not be trusted even though a particular value seems to be within its valid domain.
- Transitory checking (e.g., close time after boot sequence) necessary to fulfil plausibility check when sensors are not yet available. Generally, “pending” status flag is set.

Since a substitute value is used it may be useful to notify the receiver that the original value is not available using a signal qualifier. The receiver can then decide on how to interpret and use the value.

Note that there are no generic signal qualifier mechanisms in AUTOSAR. Applications should define their own mechanism, for example by transmitting a record with a value and a qualifier.

8.2.2 Applicability

Table 8-3: Applicability matrix for *substitute values*.

Step\Error Model	Data	Program flow	Access	Timing	Asymmetric
Detection					
Isolation					
Recovery	X	X		X	X

Assigning a substitute provides means for subsequent operations to proceed in a way that makes an end result useful. However, it will not provide recovery from an error as it is not alleviating the situation that lead to the erroneous state in first place. Thus, it is partially applicable to recovery as it allows an existing error to be masked to a certain degree.

As the erroneous state that is to be masked with substitute values can be the result of any type of underlying error, it is partially applicable to all types.

8.2.3 Application level vs. BSW

Substitute values often require application knowledge that is not present at platform level and assignment is in these cases performed in the SW-Cs. At configuration time BSW modules can be configured with default substitute values. However, application specific knowledge is still typically needed.

8.2.4 AUTOSAR References

Table 8-4: AUTOSAR References for *substitute values*.

Name	Type	Document	Comment
Software Component Template, RTE, AUTOSAR COM	SWS Template	SWC-T [1] RTE [2] & COM [4]	Usage of initial and default values.
Software Component Template, NVM	SWS Template	SWC-T [1] NVM [11]	Usage of default ROM block or redundant block.

The SW-C designer can specify an `initValue` on the `UnqueuedReceiverComSpec`. This value is used when no values were received, but the application reads the value. It can also be used in case of invalid values, depending on the `UnqueuedReceiverComSpec`'s `handleInvalid` attribute (`dontInvalidate`, `keep`, `replace`). These `initValues` are implemented by COM ([4]) or RTE ([2]) based on the SW-C XML description (see the Software Component Template [1]).

An `initValue` can also be specified for `UnqueuedSenderComSpec`, `InterRunnableVariables`, `PerInstanceMemory`, or the `ramBlocks` of `NvBlockComponentTypes` (see the Software Component Template [1]).

The NVM ([11]) can use a default ROM block in case of failure. This block can be defined by the SW-C designer with a `ParameterDataPrototype` in the `defaultData` role (see `RoleBasedDataAssignment` in the Software Component Template [1]).

8.3 Voting

8.3.1 Description

A basic principle for building fault-tolerant systems is to execute fragments ("components") redundantly and then consolidate the results of each component by performing a vote on the results. The actual vote is typically performed by a dedicated component called "voter". Common voting algorithms include "simple majority", "2 out-of 3" etc.

Voting can be performed at multiple levels, from replicated runnables in one SW-C to application level voting across ECUs. Replication can be made on the binary/source code level or on the specification level. In the former case each component is a copy

of the same original component, whereas in the latter case components are different, but are built using the same specification.

Contrary to the comparison mechanism presented in Section 8.1, voting can handle more than two replicas and, depending on the number of replicas voting also provide isolation (identification of faulty replica) and partial recovery (a good value is output).

8.3.2 Applicability

Table 8-5: Applicability matrix for *voting*.

Step\Error Model	Data	Program flow	Access	Timing	Asymmetric
Detection	X				
Isolation	X				
Recovery	X				

Voting is used to detect data errors. Furthermore, the source of the erroneous value can be identified if three or more values are voted upon (and at most one value is erroneous). As a correct value can be produced despite the presence of an error, voting partially supports recovery. Additional means are needed to fully recover from the error.

8.3.3 Application level vs. BSW

Voting is performed in SW-Cs.

8.3.4 AUTOSAR References

AUTOSAR does not provide a voting service. Since voting mechanisms are not provided by the AUTOSAR BSW, these need to be implemented at application level in the SW-Cs who require them. AUTOSAR supports multiple instantiation of SW-Cs. This feature can be used by a SW-C implementer to implement a specific voting mechanism for an application.

8.4 Agreement

8.4.1 Description

When redundant components are used to increase the reliability of an application agreement may be needed for components (called participants) to *agree* on the value

used (including the result of some computation) by exchanging the result of local computations as messages².

The difference between agreement and voting mechanisms is that when using agreement components interact to reach a decision, whereas in voting it is left to the voter to decide. Agreement protocols could be compared with closed loop systems, where the feedback consists of the sent messages received by all other parties. Voting can analogously be compared to an open loop system where the voter collects values and decides.

Agreement protocols can also handle asymmetric faults through multiple rounds of information exchange. Thus all (correct) participants agree on the same value as well as on the correctness of the other participants.

8.4.2 Applicability

Table 8-6: Applicability matrix for *agreement*.

Step\Error Model	Data	Program flow	Access	Timing	Asymmetric
Detection	X				X
Isolation	X				X
Recovery	X				X

As values are compared similar to voting, agreement mechanisms can detect and isolate data value errors. The extra rounds of information exchange allow also for detection and isolation of asymmetric errors. Agreement cannot handle either program flow or timing errors.

Recovery is partially supported in that a faulty participant can be identified and its behavior masked from affecting the system. Additional means are needed to fully recover from the error.

8.4.3 Application level vs. BSW

Basic services are implemented on BSW-level, but applications using agreement protocol must be aware of the fact they participate. Proposal of new values and adoption of agreed values are examples of situations where applications need to be aware of the protocol.

² Obviously other communication paradigms can be used, messages are used as illustration.

8.4.4 AUTOSAR References

No true agreement service exists in AUTOSAR. If specific agreement semantics are needed for application level communication it shall be implemented specifically for these applications.

8.5 Checksums/Codes

8.5.1 Description

A technique for increased data consistency is to add redundant information to the data values to protect. The extra information allows for detection of modifications of (parts of) the data, and in some cases even correction and restoration of the original data values. The cost is both in terms of performance (time to calculate and check the checksum, additional communication needs) and in terms of the additional memory requirements for storage. Extensions also include cryptographic algorithms providing digital signatures and encryption/decryption of data.

There are multiple uses of checksums/codes, including:

- Safely storing data in both volatile and non-volatile memory
- Dependable communication between SW-C, both inter- and intra- ECU
- Protecting data from tampering (data integrity) by unauthorized entities
- Sending and receiving of encrypted data across unsafe channels

Please note the difference between the first two cases, which are concerned with benign errors, whereas the latter two are concerned with malicious errors (attackers, intruders etc). However, the same mechanisms can often be used for multiple purposes.

Additional threats to data security include spoofing (pretending to be someone else), repudiation (denying a performed action), denial of service, and elevation of privileges. In general this document is focused on benign errors. However, these threats may be of importance for some applications.

8.5.2 Applicability

Table 8-7: Applicability matrix for *checksums/codes*.

Step\Error Model	Data	Program flow	Access	Timing	Asymmetric
Detection	X				X
Isolation	X				X

Recovery	X				X
----------	---	--	--	--	---

Checksums and codes are mainly targeted at protection of data throughout the FDIR process. Depending on the type of code used, all steps of the process can be supported.

Codes are also used as part of certain protocols (e.g., agreement) for handling asymmetric faults.

8.5.3 Application level vs. BSW

Checksums and cryptographic libraries are implemented on BSW-level or as libraries, due to performance and portability reasons. The use of dedicated peripheral circuits further decreases the use of application level mechanisms.

8.5.4 AUTOSAR References

Table 8-8: AUTOSAR References for *checksums/codes*.

Name	Type	Document	Comment
Crypto Service Manager	SWS	CSM [13] CAL [14]	Cryptographic checksum and codes
CRC Routines	SWS	CRC [12]	CRC routines
SW-C End to End Communication Protection	SWS	E2E Library [19]	Additional CRC added to signals by SW-Cs

The CSM can be used by SW-Cs to compute cryptographic checksums or codes through a port interface.

The CAL and CRC are libraries which can be used directly by SW-Cs to compute checksums.

The SW-C End-to-End communication protection library can be used to define a protocol and protect with a checksum or code the data sent by SW-C through port interfaces.

8.6 Execution sequence monitoring

8.6.1 Description

Correct execution of an application includes that the sequence of executable entities is correct. Monitoring execution sequence will enable the detection of erroneous execution paths that may result in erroneous results.

Monitoring of execution sequence can be performed at different levels of granularity. Some examples of levels of granularity are:

- Individual statements – This is the finest granularity with which execution sequence can be monitored at source code level. The sequence of individual statements in the code is monitored.
- Basic blocks – A basic block is a block of code that has exactly one entry point and one exit point and cannot be entered or exited outside of these two entry and exit points. Thus execution from entry point to exit point is strictly sequential. Note that the minimum basic block is a single statement. The execution sequence of basic blocks can be specified in a so called *control flow graph*, and monitoring execution sequence at this granularity would be to ensure that execution is performed according to this graph. Control flow in this context is synonymous to program flow.
- Runnables – A runnable has one entry point but may potentially have multiple exit points and several valid execution paths from entry point to these exit points. Also, the paths may include loops. At this level the execution sequence of the runnables of one (or more) application(s) is monitored.

Depending on the granularity of the monitor, the resource requirements may range from fairly low to very high. Monitoring the sequence of individual statements is likely to require huge amounts of memory and processing time, whereas the sequence of runnables could likely be monitored with very low overhead in memory and execution time.

8.6.2 Applicability

Table 8-9: Applicability matrix for *execution sequence monitoring*.

Step\Error Model	Data	Program flow	Access	Timing	Asymmetric
Detection		X			
Isolation					
Recovery					

Monitoring of execution sequence will detect errors in program flow. These errors may be the result of previous data errors, timing errors or asymmetric errors. However, this cannot be distinguished by the monitor itself.

A monitor can not be used to recover from an error as it only checks the current state against some predefined notion of correctness (this is true for all kinds of monitors, not only those for execution sequence).

8.6.3 Application level vs. BSW

The most practical approach is probably cooperation between applications and the BSW where the application provides the BSW with information on where it is in the execution trajectory and the BSW then checks whether this location is a valid one. This requires predefined valid trajectories. One approach could be to configure valid successors from a given location in the execution trajectory for a set of locations. The granularity (instruction, basic block, runnable) could then be defined at configuration time.

8.6.4 AUTOSAR References

Table 8-10: AUTOSAR References for *execution sequence monitoring*.

Name	Type	Document	Comment
Watchdog Manager	SWS	WdgM [10]	supervision counter and program flow monitoring
SW-C End-to-End Communication Protection	SWS	E2E Library [19]	Data sequence control on messages
AUTOSAR COM	SWS	COM [4]	Sequence counters for messages sent over the bus.

The Watchdog Manager can monitor heartbeats from application components not only for time but also for sequence. The correct sequence of execution is configured by the developers. The configuration contains the definition of a set of checkpoints or spy points and for each such point a set of allowed successors. It is then the responsibility of the SW-Cs (i.e. the developers) to make sure that each checkpoint/spy point is reported to the Watchdog Manager, which then checks the execution for temporal as well as logical sequence. In case of error detection, the ordinary recovery capabilities of the Watchdog Manager are utilized.

AUTOSAR COM and the SW-C E2E Communication Protection Library permit to check whether a stream of instances of a signal have been received in sequence.

8.7 Aliveness monitoring

8.7.1 Description

Aliveness monitoring deals with checking whether entities in a system are alive and well, i.e., are running as expected, in terms of periodicity or execution instances. A common way of monitoring aliveness is to monitor heartbeats from the parts that shall be monitored. If the heartbeat is within a certain range (minimum and maximum pulse), the monitored entity is said to be alive and well.

This mechanism is complementary to execution time monitoring, and deals with arrival rates rather than the time spent in a calculation.

Aliveness monitoring could be done at several levels:

- At application level – for example, if the application wants to monitor its internal components and thus, the various SW-Cs of the application could provide heartbeat signals.
- At BSW level – for example, the BSW could monitor that the application components on an ECU behave as specified (within the limitations of the used monitoring principle).

8.7.2 Applicability

Table 8-11: Applicability matrix for *aliveness monitoring*.

Step\Error Model	Data	Program flow	Access	Timing	Asymmetric
Detection				X	
Isolation					
Recovery					

8.7.3 Application level vs. BSW

The heartbeats would have to be generated by the applications, but the checking of the pulse could be done in the BSW. Each stream of heartbeats can be configured for a certain pulse range (minimum and maximum thresholds) and ECU mode (Heartbeats depend from the ECU mode: boot, standby, etc.).

8.7.4 AUTOSAR References

Table 8-12: AUTOSAR References for *aliveness monitoring*.

Name	Type	Document	Comment
Watchdog Manager	SWS	WdgM [10]	The hub for collecting aliveness proofs and triggering the hw watchdogs through the Watchdog interface and drivers.
AUTOSAR OS	SWS	OS [3]	Provides the possibility to restart tasks.

The watchdog manager can be used to supervise SW-Cs and/or BSW Modules. Which “entities” to supervise is pre-configured, together with the supervision parameters (no complete list, see WatchdogManager SWS [10]):

- The expected number of *aliveness indications* within a certain amount of *supervision reference cycles*.
- Tolerance levels on the detection.
- The tolerable number of *failed reference cycles*

A supervised entity can be in one of three states, WDGM_MONITORING_OK, WDGM_MONITORING_FAILED or WDGM_MONITORING_EXPIRED. When the state changes from WDGM_MONITORING_OK to WDGM_MONITORING_FAILED recovery can be initiated and if successful (the number of aliveness indications has reached the tolerable limit before the number of tolerable *failed reference cycles* is reached) the watchdog is triggered and no actions are performed. When the number of allowed *failed reference cycles* is exceeded (the monitoring of the supervised entity has failed permanently), the state changes from WDGM_MONITORING_FAILED to WDGM_MONITORING_EXPIRED and the watchdog will not be triggered anymore.

8.8 Status and Mode Management

8.8.1 Description

Status and mode management deals with meta information for signals, applications, devices, etc. This meta information can be used to analyse the state of the system in order to isolate a faulty subsystem/component and modify its behaviour accordingly. This type of information is defined at various levels:

- Signal status: In addition to the value, a signal may have meta information associated with it, such as

- Signal quality: This indicates the quality of the value, such as *nominal value*, *modelled value*, *replacement value*, and *default value*. The receiver of the signal may react differently depending on the quality of the signal.
- Signal timestamp: This indicates when the value was created and can be used to check the age of a signal. If a calculation uses several input signals, one can check that all input signals are created within a tolerable time window.
- Signal sequence number: It may be of interest to check that signal values are received in a certain order, and that no values are lost between reads. Also, sequence numbers can be used in a similar fashion as timestamps in that they allow a check that all values in a group are from a particular creation window (same sequence number, or with a minimum/maximum deviation).
- Update information: A consumer of a signal may want to know whether a signal has been updated since the last read or not.
- Device status: An application may want to know the status of the devices it uses, such as sensors and actuators. If a device is not in normal operation, applications may want to choose to deliver some form of degraded service.
- Application status/mode: An application can also have a status, or a mode, which indicates the overall health or operating situation of the application. This status/mode can be used for recovery purposes, both internally in the said application and externally by other application.
- Vehicle mode: A vehicle may be in a number of different modes (e.g. normal operation, parked, limp-home) and the applications will have to behave accordingly.
- ECU Mode: An ECU may be in different states, such as sleeping, running, powered down, and transitional states between such states.

In order to provide support for this kind of status and mode management, it must be possible to set and get this information at application level (although there may only be one producer of a particular piece of status information there may be many consumers).

8.8.2 Applicability

Table 8-13: Applicability matrix for *status and mode management*.

Step\Error Model	Data	Program flow	Access	Timing	Asymmetric
Detection					
Isolation	X	X		X	X

Recovery					
----------	--	--	--	--	--

From an error handling point-of-view, status and mode information can be used in the isolation phase. We have chosen to tag this mechanism as applicable for recovery as this mechanism in itself is not useable for recovery. Recovery can of course be triggered as a result of the status or mode on the monitored entity. However, the action of recovering itself is not part of this mechanism.

8.8.3 Application level vs. BSW

BSW modules manage meta information and distribute this information from producers to consumers. The definition, setting and getting of this information is done at application level though.

8.8.4 AUTOSAR References

Table 8-14: AUTOSAR References for *status and mode management*.

Name	Type	Document	Comment
ComM	SWS	ComM [5]	ComM handles the communication modes of the ECU and can trigger a shutdown of the bus if no communication is required. It also implicitly keeps the ECU alive (interactions with EcuM-Fixed or BswM).
RTE	SWS	RTE [2]	RTE Spec, including application mode management Communication of modes
Software Component Template	SWS	SWC-T [1]	Application modes defined Modelling of the modes communication
BSW Mode Manager	SWS	BswM [15]	Modes and transitions management.
ECU State Manager	SWS	EcuM [7]	The ECU State manager manages the state of a single ECU.

Operating System	SWS	OS [3]	States of OS Applications. OS has states for other OS objects (tasks, resources, etc.) but those are not directly accessed by applications.
------------------	-----	--------	--

The RTE defines status for read data values, through the `Rte_Read`, `Rte_IStatus`. Status can be `RTE_E_OK`, `RTE_E_INVALID` or `RTE_E_MAX_AGE_EXCEEDED`. `RTE_E_INVALID` refers to an explicitly invalidated data value and `RTE_E_MAX_AGE_EXCEEDED` refers to an outdated data element.

The RTE allows a SW-C to specify mode, which can be use to execute or inhibit runnables. Modes can be defined for application specific purposes. This should be sufficient for error handling purposes as well.

The ECU State manager (and the Basic Software Mode Manager when it is used) manages the state of a single ECU.

8.9 Reconfiguration

8.9.1 Description

A technique for building fault-tolerant systems is to detect and isolate faults and then reconfigure the system to no longer use the faulty component, or to reconfigure to provide only a degraded set of services (or level of service).

Examples of reconfiguration strategies

- Isolating faulty components by hindering further communication. This could also include shutting down components selectively.
- Reconfiguration of protocol parameters, for instance voting algorithms, tolerance levels etc.
- Degraded functionality, such as providing only ABS and no ESP or a special “limp home” mode.

Reconfiguration is typically controlled using static policies, which are configured at system configuration time. The policies define when a reconfiguration is triggered, and how it is performed. Common triggers include error signals, as is the case for the Function Inhibition Manager (FIM) defined in AUTOSAR, which is triggered by messages from the DEM upon error. The FIM is limited to only informing an application of a request to inhibit parts of SW-Cs (so called “functionalities”) and cannot actively inhibit anything or trigger a reconfiguration.

Note that there is a difference between mode management (Section 8.8: Status and Mode Management) and reconfiguration. Mode Management is an infrastructure to

transfer information on states (i.e., modes) in the system, such that certain actions can be taken. These actions may be reconfiguration actions.

8.9.2 Applicability

Table 8-15: Applicability matrix for *reconfiguration*.

Step\Error Model	Data	Program flow	Access	Timing	Asymmetric
Detection					
Isolation					
Recovery	X	X	X	X	X

Reconfiguration of SW-Cs is part of the recovery step and cannot aid in either detection or isolation. It can possibly apply to any error, given that a reconfiguration policy is defined.

8.9.3 Application level vs. BSW

Reconfiguration can be performed at the application level and with support by the BSW services. Reconfiguration due to errors may require pure BSW support (such as Reset, Section 8.10).

8.9.4 AUTOSAR References

Table 8-16: AUTOSAR References for *reconfiguration*.

Name	Type	Document	Comment
Function Inhibition Manager	SWS	FIM [8]	The FIM can be used to selectively deactivate SW-C functionalities.
BSW Mode Manager	SWS	BswM [15]	The BswM arbitrates mode requests from SW-Cs in the application layer and performs mode switches based on pre-defined rules.
Terminate or Restart Partitions	Implementation note	Section 10	Defines requirements for terminating and restarting partitions, which is a feature involving multiple BSW Modules.
AUTOSAR OS	SWS	OS [3]	The OS handles termination of OS-Applications.
RTE and Software Component Template	SWS Template	RTE [2] SWC-T [1]	Definition and support of modes. Definition and support for termination and restart.

The FIM provides a control mechanism for SW-Cs and the functionality therein. In this context, a functionality can be built up of the contents of one, several or parts of runnable entities with the same set of permission / inhibit conditions. By means of the FIM, inhibiting (→ deactivation of application function) these functionalities can be configured and during runtime facilitating reconfiguration of the application.

The BswM can be configured to switch mode of the BSW based on mode requests. The interaction with SW-Cs is performed through the RTE using ModeDeclarationGroups. The BswM thus performs two basic tasks: Mode Arbitration and Mode Control. The Mode Arbitration part initiates mode switches resulting from rule based arbitration of mode requests and mode indications received from SW-Cs or other BSW modules. The Mode Control part performs the mode switches by execution of action lists containing mode switch operations of other Basic Software modules. The action lists associated with a mode switch can be used to reconfigure the application, such as start/stop of I-PDUs (COM), disable all communication (NM) or changed PduR routing etc.

Termination/reconfiguration using the FIM requires cooperation with the affected SW-Cs implying correct behaviour from the SW-C. In case of malfunctioning SW-Cs a more brute force approach is required which terminates SW-Cs independently of the correctness of their behaviour. Using an approach as described in Chapter 10, SW-Cs belonging to a partition can be forcibly terminated.

The AUTOSAR OS handles termination of OS-Applications, to which tasks and other OS resources belong. The OS provides a service to terminate an OS-Application (TerminateApplication). This OS service can be called from an SW-C belonging to a trusted OS-Application.

8.10 Reset

8.10.1 Description

An application may try to recover in a number of ways, ranging from setting replacement values and wait, hoping that the error will disappear by itself (transient/intermittent errors) to full reconfiguration of its structure in order to shutdown faulty components and launch backups and replacement components (permanent faults). Sometimes, it may not be possible to do these things though and a complete reset is necessary to start from a known good state. This holds not only for transient HW faults, but also for “soft” SW faults (systematic faults), sometimes referred to as Heisenbugs [2]. Such faults are transient in nature and cannot be easily repeated. A reset of the SW state typically removes Heisenbugs, as it puts the SW back in a known and well tested state.

Resets are used to recovery from transient faults. Permanent faults (e.g., permanently defect HW or SW bugs) cannot be recovered from. Other mechanisms to isolate transient faults are therefore needed to avoid using resets for permanent faults.

Reset can potentially be performed at the following levels:

- SW-C reset. An SW-C is found to be faulty and is reset in order to get it back into a safe state. The reset takes place at the application level.
- Application reset. If it is not sufficient to just reset single SW-Cs, it may be necessary to restart the whole application so that it can resume its normal service. The reset affects several SW-Cs at the application level, and may involve SW-Cs at multiple ECUs.
- ECU reset. If all else fails, it may be necessary to reset the entire ECU on which the fault or error has been found. This kind of reset will affect all applications that have SW-Cs located on the ECU as well as the BSW. The reset will also likely be visible to other ECU's on the network.

8.10.2 Applicability

Table 8-17: Applicability matrix for *reset*.

Step\Error Model	Data	Program flow	Access	Timing	Asymmetric
Detection					
Isolation					
Recovery	X	X	X	X	X

A reset is a last resort to recovery and is only applicable in that phase.

8.10.3 Application level vs. BSW

Reset at SW-C and application level could be controlled at application level, i.e., a dedicated SW-C could detect that an application has failed and request a reset. The RTE/OS shall provide services that enable the distribution of reset commands to the affected SW-Cs. For ECU reset, the BSW must be responsible for performing this. In the ECU case, all affected applications could be made aware of the impending reset in order to prepare themselves for it.

Note that even though the reset is initiated at the application level it will always require RTE/OS support for performing it.

8.10.4 AUTOSAR References

Table 8-18: AUTOSAR references for *reset*.

Name	Type	Document	Comment
Watchdog Manager	SWS	WdgM [10]	The hub for collecting aliveness proofs and triggering the hw watchdogs through the Watchdog interface and drivers.
AUTOSAR OS	SWS	OS [3]	Provides the possibility to terminate and restart OS-Applications. Controls and monitors timing behaviour of tasks.

ECU State Manager	SWS	EcuM [7]	An SW-C can select the shutdown target, i.e., which activity shall be performed after an EcuM controlled shutdown.
Terminate or Restart Partitions	Implementation Note	Section 10	Defines requirements for terminating and restarting partitions, which is a feature involving multiple BSW Modules.
Diagnostic Communication Manager	SWS	DCM [6]	Diagnostic Communication Manager can initiate resets through MCU

The ECU State Manager provides an interface to SW-Cs for selecting different shutdown targets, that is, what the ECU shall do when a shutdown is performed by the EcuM. The shutdown target can be either sleep, reset or off.

The WdgM monitors SW-Cs based on the aliveness indications made. The WdgM provides three mechanisms for the aliveness monitoring: a) supervised entities, b) temporal program flow monitoring, and c) logical program flow monitoring [10]. Missing aliveness proofs from applications can trigger a variety of actions by the WdgM:

- Inform the offending SW-C using the mode management mechanism in the RTE (“Local Failure Recovery”).
- Inform the DEM. SW-Cs (the offending one and/or others) can then react upon DEM notifications. (“Global Failure Recovery”).
- Termination or restart of a partition using the TerminateApplication service provided by the OS.
- Indicate to the watchdog driver that it shall cease triggering the HW watchdog, eventually leading to an ECU reset.
- Directly resetting the ECU through the MCU.

The AUTOSAR OS provides the possibility to terminate or restart an OS-Application, which is a set of OS resources (including tasks, i.e., SW-Cs). This can be triggered either by a protection violation (such as memory or timing) or manually by a trusted SW-C. The manual reset request from a SW-C makes it possible to reset even distributed applications in a coordinated fashion. When the OS-Application is restarted a dedicated restart task is used to perform restart activities and is responsible for notifying involved BSW Modules of the restart.

The DCM can perform an ECU reset upon a diagnostic request from an external diagnostic client (Tester). Such clients could potentially be vehicle-internal and would then be able to request ECU resets based on the observed state of the vehicle. However, this is marginally at application level and will not be considered further here.

8.11 Error Filtering

8.11.1 Description

In some situations taking recovery actions due to errors, for instance transients, may cause more damage than it does good. Reacting to such errors may cause an over-reaction, where the recovery actions may put the system in a state where it is less safe than previously (for instance while restarting ECUs). In such cases a filtering of the errors may be needed before certain recovery actions are taken.

A common example is discrimination between transient and permanent errors using counters, where erroneous behavior increases the counter and correct behavior decreases it. When it reaches a specific threshold, the error is classified as permanent (a failure) and recovery is initiated.

8.11.2 Applicability

Table 8-19: Applicability matrix for *filtering*.

Step\Error Model	Data	Program flow	Access	Timing	Asymmetric
Detection					
Isolation	X	X		X	X
Recovery					

Error filtering is only applicable as a means for isolation. It requires additional detection mechanisms and can trigger recovery mechanisms. However, filtering of errors may alleviate the need for unnecessary recovery actions, and thus contribute also to a better recovery strategy by gaining information on the nature of the error.

8.11.3 Application level vs. BSW

A central debouncing mechanism is provided by the DEM within BSW. On the application level, application-specific error filtering can be applied, e.g., classifying transient errors. Such classification needs to be implemented by the SW-Cs.

8.11.4 AUTOSAR References

Table 8-20: AUTOSAR references for *filtering*.

Name	Type	Document	Comment
------	------	----------	---------

Specification of Diagnostics Event Manager	SWS	DEM [9]	The DEM provides de-bouncing mechanisms to confirm errors.
--	-----	---------	--

8.12 Memory Protection

8.12.1 Description

Memory protection is used to protect against errors propagating from one protection domain (partition) to another. Partitions are defined to form error confinement regions, where applications can be placed for mutual protection. Such protection enables separation between applications and thus enables multiple suppliers of SW-Cs to deliver SW for an ECU. This is important both for analyzing and enforcing safety issues.

8.12.2 Applicability

Table 8-21: Applicability matrix for *memory protection*.

Step\Error Model	Data	Program flow	Access	Timing	Asymmetric
Detection			X		
Isolation ³					
Recovery			X		

Memory protection is mainly an error detection mechanism (for memory access errors). However, as the execution is halted before the write has been performed error propagation is confined. To achieve full recovery additional mechanisms are required, for instance to terminate or restart the execution of the offending runnable/task/SW-C/partition without endangering the execution of other applications (See Section 8.9 - Reconfiguration). Therefore memory protection only partially supports recovery from memory access violations.

8.12.3 Application level vs. BSW

Memory protection mechanisms are implemented in the BSW with HW support. To fully support recovery the RTE needs to also be aware of the protection mechanisms and act when applications are terminated or restarted. The application needs not be

³ See the definition of isolation in section 5.2. A memory protection mechanism can be used to stop the propagation of an error, but it cannot identify the source of the error.

aware that it is running in a specific partition, but if restart is enabled it must be built to handle consistency issues arising from differing internal states as a result of a restart.

Memory protection is not a mechanism used directly by SW-Cs, but it is configured by application developers/ECU integrators and is therefore relevant for application developers.

8.12.4 AUTOSAR References

Table 8-22: AUTOSAR references for *memory protection*.

Name	Type	Document	Comment
AUTOSAR OS	SWS	OS [3]	Provides basic memory protection, the possibility to terminate and restart OS-Applications, and is involved in communication across protection boundaries.
RTE	SWS	RTE [2]	The RTE is involved in the termination and restart of partitions, and is involved in communication across protection boundaries by ensuring communication consistency.
Terminate or Restart Partitions	Implementation Note	Section 10	Defines requirements for terminating and restarting partitions, which is a feature involving multiple BSW Modules.

The AUTOSAR OS together with the RTE implements the memory protection facilities in the system. The OS provides the fundamental protection mechanisms together with HW support and the OS and the RTE facilitate the communication across protection boundaries.

For an application developer it is important to know if the application is to be put in a partition for protection and what actions will be taken in case of protection violations. In case of restarts of a partition the applications must handle any inconsistencies that may arise when partitions are terminated or restarted. The system only guarantees consistency of communication while restarting and that init runnables will be executed during restart. The developer must therefore cooperate with the system integrator, which is responsible for the restart code, to find suitable solutions.

8.13 Timing Protection

8.13.1 Description

Timing protection refers to protecting the system against activities requiring too much time to complete, such as an executing component taking too much execution time on the processor and thereby hindering the execution of other components, communication delays, peripheral units not responding in time, etc.

For activities, one may define time budgets providing an upper limit on how much time a given activity may use. For example, one may choose to set execution time budgets for components, or a maximum response times for communication.

8.13.2 Applicability

Table 8-23: Applicability matrix for *timing protection*.

Step\Error Model	Data	Program flow	Access	Timing	Asymmetric
Detection				X	
Isolation					
Recovery				X	

Execution time monitoring can be used to detect that a SW-C (actually a task) has exceeded its assigned execution time budget. It cannot detect timing errors in communication directly. Communication time-out monitoring can detect when a response is not received within the expected time.

To fully recover additional mechanisms are needed (reset, reconfiguration, etc.) and thus timing protection only partially supports recovery.

8.13.3 Application level vs. BSW

Timing protection is implemented in the BSW, i.e., the BSW performs the actual monitoring. Violations may be reported to SW-Cs. Execution time budgets and communication response deadlines are configured by application developers/ECU integrators and are therefore relevant for application developers.

8.13.4 AUTOSAR References

Table 8-24: AUTOSAR references for *timing protection*.

Name	Type	Document	Comment
AUTOSAR OS	SWS	OS [3]	Provides the possibility to terminate and restart OS-Applications. Monitors timing behaviour of tasks.
RTE	SWS	RTE [2]	The RTE is involved in the termination and restart of partitions and ensures platform consistency. Timeout monitoring
Terminate or Restart Partitions	Implementation Note	Section 10	Defines requirements for terminating and restarting partitions, which is a feature involving multiple BSW Modules.
AUTOSAR COM	SWS	COM [4]	COM provides deadline monitoring for signals.
Watchdog Manager	SWS	WdgM [10]	The hub for collecting aliveness proofs and triggering the hw watchdogs through the Watchdog interface and drivers.
Software Component Template	Template	SWC-T [1]	The SWC-T defines the requirements for the timeout handling, whether a SW-C supports restart, the Watchdog service needs

The AUTOSAR OS provides basic timing protection facilities to monitor execution of tasks and ISRs. When a timing violation occurs the ECU-wide protection hook is called which has the possibility to terminate tasks or OS-Applications, shut down the OS or do nothing. As the scope for timing violation reactions can also be OS-Application wide, one can consider these mechanisms to act on partitions, and the developer/integrator can partition the system accordingly.

For an application developer it is important to know if the application is to be put in a partition for protection and what actions will be taken in case of protection violations. In case of restarts of a partition the applications must handle any inconsistencies that may arise when partitions are terminated or restarted. The system only guarantees consistency of communication while restarting and that init runnables will be

executed during restart. The developer must therefore cooperate with the system integrator, which is responsible for the restart code, to find suitable solutions.

Timeouts (called `aliveTimeout` in RTE [2]) can be defined for data elements exchanged using the RTE using sender-receiver communication. For communication on busses these correspond to signals for COM, which provides deadline monitoring, both for reception and transmission. Similarly, the RTE provides communication timeout monitoring for client-server communication. These types of deadline monitoring can be used to (for instance) detect that SW-Cs residing in terminated partitions no longer execute.

The WdgM monitors SW-Cs based on the aliveness indications made. The WdgM provides three mechanisms for the aliveness monitoring: a) supervised entities, b) temporal program flow monitoring, and c) logical program flow monitoring. See Section 8.10 and [10] for more information.

9 Aspect mapping

Each error handling mechanism is characterized by a number of properties, such as where in the FDIR process it applies or which error models it can handle. To give an overview of these different aspects of the different mechanism this section presents a number of mapping tables, where each mechanism is mapped onto the different properties. Each table gives references to the mechanisms covered individually in Section 8.

9.1 Mapping to FDIR process and Error Model

Not all mechanisms can be used for all steps in the FDIR process, and similarly they are applicable only for specific error models. To illustrate these 3-dimensional relationships (Mechanism x FDIR step x Error Model) we present three tables in this section. Each table shows an overview of the mechanisms and their applicability in each respective step of the FDIR process with respect to each of the error types defined in the error model.

More information regarding the capabilities of each error handling mechanism is found in Section 8, where each mechanism is presented in more detail.

In Table 9-1, the mechanisms are mapped to the first step of the FDIR process – detection, Table 9-2 contains the mapping to the second step – isolation, and Table 9-3 shows the mapping to the third and last step – recovery.

Table 9-1: Mapping of mechanisms to the steps of the FDIR process. Error detection.

Mechanism	Ref	Data	Program flow	Access	Timing	Asym metric
Plausibility checks	8.1	X				
Substitute values	8.2					
Voting	8.3	X				
Agreement	8.4	X				X
Checksums/Codes	8.5	X				
Execution sequence monitoring	8.6		X			
Aliveness monitoring	8.7				X	
Status & Mode Management	8.8					
Reconfiguration	8.9					
Reset	8.10					
Error filtering	8.11					
Memory protection	8.12			X		
Timing protection	8.13				X	

Table 9-2: Mapping of mechanisms to the steps of the FDIR process. Error isolation.

Mechanism	Ref	Data	Program flow	Access	Timing	Asym metric
Plausibility checks	8.1	X				
Substitute values	8.2					
Voting	8.3	X				
Agreement	8.4	X				X
Checksums/Codes	8.5					
Execution sequence monitoring	8.6					
Aliveness monitoring	8.7					
Status & Mode Management	8.8	X	X		X	X
Reconfiguration	8.9					
Reset	8.10					
Error Filtering	8.11	X	X		X	X
Memory protection	8.12					
Timing protection	8.13					

For error isolation it is important to note that the explicit information gained by detecting the error is not considered in Table 9-3. For example, when detecting that an entity has crashed by some aliveness monitoring mechanism the crashed entity is explicitly identified, however, no additional information is gained from the monitoring mechanism that can help in recovery, like the underlying reason for the crash. This is in contrast to for instance agreement on a data value, where not only the error is detected (some participant is faulty), but additionally also which participant.

The main purpose of status and mode management is the spread of error information, making it available to interested parties, and thereby making error isolation possible.

Table 9-3: Mapping of mechanisms to the steps of the FDIR process. Error recovery.

Mechanism	Ref	Data	Program flow	Access	Timing	Asymmetric
Plausibility checks	8.1					
Substitute values	8.2	X	X		X	X
Voting	8.3	X				
Agreement	8.4	X				X
Checksums/Codes	8.5	X				
Execution sequence monitoring	8.6					
Aliveness monitoring	8.7					
Status & Mode Management	8.8					
Reconfiguration	8.9	X	X	X	X	X
Reset	8.10	X	X	X	X	X
Error Filtering	8.11					
Memory protection	8.12			X		
Timing protection	8.13				X	

9.2 Mapping to implementation level

The implementation level refers to the level where the mechanism is most suitably implemented. However, the use and control of the mechanism is still in SW-Cs (or at least by application developers), i.e., on the application level.

Two implementation levels are relevant for the presented mechanisms, application level (SW-C) and basic SW level (BSW). The fundamental difference lies in where the mechanism is implemented, as a “service” provided to SW-Cs from the BSW level or as a pure application level mechanism, not requiring any specific BSW support. HW-based solutions are included at the BSW-level since direct access to HW is generally not permitted. As an example, dedicated cryptographic peripherals could be accessed through the same BSW interface as SW-based solutions.

Table 9-4: Mapping of mechanisms to implementation level

Mechanism	Reference	Implementation level
Plausibility checks	8.1	SW-C
Substitute values	8.2	SW-C / BSW
Voting	8.3	SW-C
Agreement	8.4	SW-C
Checksums/Codes	8.5	SW-C / BSW
Execution sequence monitoring	8.6	SW-C / BSW
Aliveness monitoring	8.7	SW-C / BSW
Status & Mode Management	8.8	SW-C / BSW
Reconfiguration	8.9	SW-C / BSW
Reset	8.10	SW-C / BSW
Error Filtering	8.11	SW-C / BSW
Memory protection	8.12	BSW
Timing protection	8.13	BSW

10 Terminating and restarting partitions

This section presents an approach for enabling the termination and restarting of the execution of partitions as an error handling alternative to cooperative recovery or ECU reset.

10.1 Introduction

A partition denotes the abstract notion of a logical grouping of SW-Cs and BSW resources and defines an error containment region, whereas the term OS-Application is used as the basis for the implementation of such a partition in the AUTOSAR OS. OS-Applications contain tasks, ISRs, OS Alarms etc. that are grouped together and are handled by the OS independently. As an OS-Application is a strictly OS-level notion, the existence of these is not known to SW-Cs and other BSW Modules (Basic Software Modules).

The main focus is on handling faults which lead to one of the following:

- Protection violation caused by erroneous memory access. This is detected by the OS (with support from an MMU).
- Protection violation caused by exceeding of allocated execution time budget. This is detected by the OS.
- Detection of application states which do not trigger a protection violation but require termination or restart of a partition.

Terminating or restarting a partition is achieved by terminating or restarting the underlying OS-Applications and performing any BSW and other cleanup necessary to set related resources into a consistent state. OS-Applications can be terminated and restarted independently, thereby forming error containment regions at the OS-level. Terminating or restarting a partition of course means terminating or restarting the SW-Cs contained in that partition. Being based on OS-Application, a partition is fully mapped to one ECU (i.e. cannot be spread over several ECUs).

This concept describes how terminating/restarting is performed, which modules are involved and how consistency is ensured. The basic principle is that communication across partitions behaves similar to inter-ECU communication, even when a partition is terminated or restarted. That is, there should be no difference in the semantics or behaviour when interacting partitions are mapped to the same physical ECU or different physical ECUs. The concept is intended for non-trusted OS-Applications containing SW-Cs only. In particular, this means that the partition(s) containing the BSW cannot be terminated or restarted. Moreover, the functionality is optional and only provides basic support for termination/restart activities. To get proper termination and restart of partitions working in a system, additional work is required from the developer/integrator (e.g. configuration, supporting code, definition of cleanup activities, etc.).

The approach can be used when no assumptions can be made on the correctness of the error handling capabilities of the application software. That is, no support is required from the application software. In fact, the application software in a partition is not even executing during termination and restart.

Note that the approach can potentially handle external transient faults as well as internal design faults. The only requirement is that the faults cause one of the failure scenarios mentioned above.

This document describes the basic support provided in the BSW, lists the responsibilities of the developer/integrator for successfully employing the functionality and provides suggestions on solution approaches.

10.1.1 Automotive Applications

An application in an automotive system built using AUTOSAR can consist of several interacting SW-Cs mapped over potentially several partitions. There may be mapping decision affecting the locations of SW-Cs/partitions, e.g. to save resources, or due to mapping constraints (such as sensors/actuators), or to provide fault-tolerance through redundancy and/or separation. Consequently, an application consists of a set of potentially distributed SW-Cs/partitions.

The mapping of which SW-Cs belong to an application or partitions is not explicitly modelled in AUTOSAR. However, the mapping is known to the system designer, who is also responsible for specifying how errors should be handled. It is therefore assumed that the designer can define whether terminating or restarting a partition is an appropriate response to errors for a specific application and if so, which partitions (and thereby which SW-Cs) can be terminated or restarted. This information is then used by the OEM to build an overall recovery policy for the vehicle. It is important to note that AUTOSAR SW-Cs cannot be generally restarted without the designer being aware of this fact and designing the application accordingly.

10.1.2 Software Partitioning & Error Containment Regions

Software partitioning is one of the chief goals for a general-purpose OS and is seeing increased attention in embedded systems as the complexity of embedded software increases. The AUTOSAR glossary defines partitioning as: “Decomposition, the separation of the whole system into functional units and further into software components”. In the context of error handling and safety the term partitioning is enhanced to also mean the definition and enforcement of *error containment regions*, i.e., the system (most of the time the OS) ensures that errors of certain types are contained in the (predefined) regions where they occur and will not propagate outside causing new errors in other regions. This includes, among other things, protecting the memory of one error containment region from illegal writes from other error containment regions and preventing application software from monopolizing the CPU. Also, enforcing strict access rules to system resources, such as system services and peripheral devices, from applications is an important issue.

When the software in a partition violates the protection rules set up by the OS its execution is typically forcibly halted and the user can choose to restart the offending software with the hope that the fault causing the violation was transient and will not be present when the application is restarted, or in case the error is deemed non-recoverable terminate the offending software in order to avoid causing interference in other partitions.

For automotive systems, error recovery is typically performed by either having the application recover using dedicated recovery mechanisms, here termed *cooperative recovery* as it requires a partially functional application to invoke the recovery mechanisms, or to reset the ECU on which the violation occurred. The first approach is effective when the fault causing the failure does not constitute a hindrance for the code of the recovery mechanism. In the AUTOSAR case, this could be implemented using, e.g., the functionality provided by FIM, DEM or BswM to modify the behaviour of SW-Cs. However, faults which put the software in a state not envisioned by the designer can generally *not* be handled by such an approach, and an ECU reset is instead used. Such resets are benign if they can be performed fast enough to not violate application deadlines and do not adversely affect other applications. A solution to the latter problem is to not co-locate applications which may, in case of ECU reset, adversely affect each other. However, in order to reduce manufacturing time, complexity and weight, it is desired to increase the level of integration, where multiple applications, possibly independent and with mixed criticality, share a single ECU. In such systems software partitioning is needed to ensure that error propagation is minimized across applications and independence in case of failures is upheld.

10.2 Rationale – Use Cases

Terminating and restarting an application is a basic approach to minimizing error propagation and performing error recovery (mainly the restart functionality). Two main use cases have been identified for AUTOSAR error handling and are presented in this section.

10.2.1 Use Case 1: Software Partitioning [UC1]

Software partitioning is used to prevent the effects of faults in one partition from spreading to other partitions on the same ECU, by specifying well defined boundaries in the space and time dimensions for the software of the partitions to execute within. By terminating a partition, so called *fail silent* [17] behaviour can be implemented. Additionally, this is an important aspect for building safety cases and for liability considerations when mixed-criticality applications execute on one ECU. See also the OS SWS [3] for more information. Commonly used examples of such techniques are memory partitioning and execution time (deadline) monitoring.

Memory partitioning and execution time monitoring can be used to detect memory access errors and timing errors, and to prevent them from corrupting software in other partitions. However, when such errors occur, a reaction from the system is necessary. The most straightforward approach is to not allow the offending software to continue its execution, i.e., to forcibly terminate it. This will prevent errors from propagating but it may put the system in an inconsistent state, since the software in the partition may have been interacting with software in other partitions and/or BSW at the time of termination. *The behaviour of the system and its consistency rules must therefore be well specified*, such that a partition can continue to execute without inadvertent disturbance from a sudden termination or restart of another partition.

AUTOSAR supports memory partitioning and execution time monitoring. It is possible to define time boundaries which a SW-C must execute within, and memory partitions which cannot be overridden, thus protecting SW-Cs in one partition from erroneous memory accesses SW-Cs in other partitions.

Both concepts use an OS-level feature called OS-Application [3]. For OS-Applications, the OS provides timing and memory protection, given that they are configured and HW support (for memory protection) is available. When an error (a protection violation) is detected, a proper reaction is called for.

10.2.2 Use Case 2: Application-level Error Handling [UC2]

Application level error handling can be used for two main purposes:

- Error handling of distributed automotive applications
- OEM-specific error handling

10.2.2.1 Error handling of distributed automotive applications

In general, an automotive application can (as discussed in Section 10.1.1) be divided into several partitions which form different error containment regions. These partitions could possibly be mapped to different ECUs.

The management of this distribution of application parts needs to be handled at application level, as the information of division into partitions and mapping of partitions to ECUs is not available at BSW level.

An application-level error manager (ALEM) would deal with issues that arise from such distribution. For example, if the termination of a partition leads to a need to terminate other partitions (not necessarily on the same ECU), this would be triggered by the ALEM.

10.2.2.2 OEM-specific error handling

Errors affecting a partition do not necessarily lead to protection violations as in UC1, but can force the partition into a state where it cannot itself recover and continue executing. It must either be terminated to minimize error propagation or restarted from a known “good” state to make progress. For such errors, the decision and action to terminate or restart a partition must be made by an external entity, since the error affecting the state of the partition could also affect its ability to recover. However, the decision on which recovery policy to use (terminate/restart, which partitions/SW-Cs etc) is OEM-specific and cannot be generalized and standardized by AUTOSAR.

10.2.2.3 Application-level Error Managers

Dealing with the two aspects described above would require the introduction of “error managers”, i.e., software applications⁴ that monitor the health of the system and initiate recovery if necessary. Figure 1 shows an example of such a setup.

Recovery may range from sending application-specific instructions, initiating reconfiguration of applications to terminating/restarting applications or the entire ECU. In this concept we focus on the termination and restarting of SW applications.

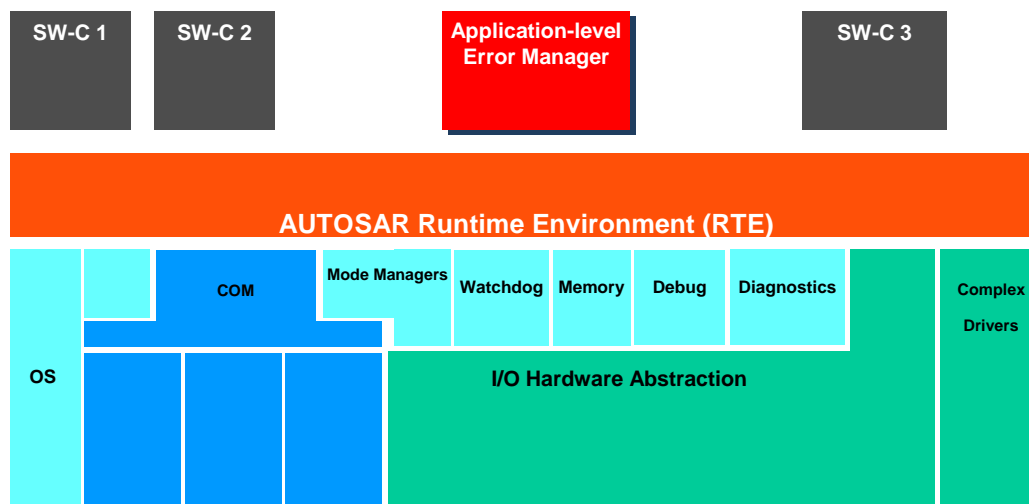


Figure 1: An Application-level Error Manager (ALEM) monitors the health of the system and initiates recovery actions according to application or OEM-specific recovery policies.

This use case is similar to the Fault Management Framework (FMF) developed in the EASIS project [25]. The proposed features in this document have corresponding functionalities defined in the FMF.

Furthermore, error management at application level would be able to take distribution into account. That is, if an automotive application is distributed over several partitions, and possibly even ECUs, and the BSW would not be able to perform a coordinated termination/restart. The case where the termination or restarting of partitions across several ECUs shall be coordinated (not necessarily synchronized) could be managed by ALEMs.

NB1: In the remainder of this section, we will refer to Application-level error manager (ALEM) whenever we mean a SW-C which is allowed to trigger partition termination and restart.

NB2: This approach does not aim to standardize an ALEM, its behaviour or its interfaces.

⁴ The ALEM could also be implemented as a BSW Module, but this is a less flexible solution. This would also entail more effort on specifying an AUTOSAR BSWM for this purpose as well as introducing the notion of a BSW Module that is potentially distributed over several ECUs. This is not within the scope of this concept.

10.3 Approach for Terminating and Restarting Partitions

In the subsequent sections, the text describes the support in AUTOSAR OS, BSW and RTE.

10.3.1 OS features

This section describes the basic functionalities and mechanisms provided by the OS that will be used to incorporate the possibility to terminate or restart a partition. Comments on how these fit with the actions necessary for termination and restart of partitions is provided here and more detailed descriptions are then described later in Section 10.3.2.

10.3.1.1 OS-Applications

The AUTOSAR OS uses a construct called “OS-Applications” (see [3]). An OS-Application is a set of OS objects, such as tasks, interrupt service routines, alarms, schedule tables, etc. From the OS point of view an OS-Application forms a cohesive functional unit. For instance application specific hook functions can be defined, such as startup and shutdown hook functions.

OS-Applications are used as a basis for partitions, i.e., protection across partitions is enforced at the OS-Application level. Software Partitioning is used for instance for memory access control or timing control, but is a general approach to address error propagation in the system. Partitions are defined to be error containment regions, i.e., with proper enforcement errors occurring in one partition cannot spread out to other partitions.

For OS controlled objects, there are a number of protection mechanisms available (see [3]). The mechanisms relevant for the termination/restart approach are memory protection and timing protection. Memory protection covers stack and data and, optionally, code of OS-Applications. Timing protection is based on execution time budgets and arrival rates.

OS-Applications are divided into i) trusted OS-Applications and ii) non-trusted OS-Applications (see [3]).

Trusted OS-Applications run in privileged mode and may execute with all (OS level) monitoring shut off. That is, a trusted application is assumed to never violate memory access rights or deadlines.

Non-trusted OS-Applications run in non-privileged mode are automatically monitored by the OS w.r.t. memory access and deadline violations. SW-Cs are allocated either to non-trusted OS-Application, or to one trusted OS-Application of BSW. There are zero, one or more SW-Cs per OS-Application.

10.3.1.2 Protection Hook

The system has a single global *Protection Hook*. The Protection Hook is called by the OS when a protection violation occurs, and its role is to decide which action is to be performed as a reaction to the error. The Protection Hook is called with a parameter value indicating the detected violation. The Protection Hook can query which OS-Application⁵ caused the violation and then decide on necessary actions the OS shall perform once the hook returns (see above). If no protection hook is configured, the OS will call `ShutdownOS()`. The return value of the protection hook defines further action performed by the OS.

- `PRO_IGNORE`: Do nothing
- `PRO_TERMINATETASK`: Forcibly terminate the faulty Task/Category 2 OsIsr
- `PRO_TERMINATEAPPL`: Forcibly terminate the faulty OS-Application
- `PRO_TERMINATEAPPL_RESTART`: Forcibly terminate the faulty OS-Application and trigger the execution of the `OSRestartTask`.
- `PRO_SHUTDOWN`: Call `ShutdownOS()`.

It is expected that the integrator, together with the application designer, decides on the error handling strategy, i.e., whether to perform a termination, restart, error filtering etc.

When an application-level request for termination or restart is issued (for details, see Section 10.3.2.4), the decision is already taken and the OS only executes the request. The Protection Hook is not executed for such an externally triggered termination or restart.

Terminating or restarting an OS-Application is an activity solely performed in the OS. In order to terminate or restart a partition, there are a set of cleanup activities which have to be performed in the BSW and potentially also in the SW-Cs of the partition. These cleanup activities will be placed in the `OSRestartTask` which is only triggered when the return code from the Protection Hook is `PRO_TERMINATEAPPL_RESTART`.

Important: For terminating or restarting a partition, the Protection Hook must return `PRO_TERMINATEAPPL_RESTART` in both cases. Otherwise there will be no opportunity to perform clean-up activities.

10.3.1.3 TerminateApplication() API

The OS provides an API for explicit termination or restart of OS-Applications. This API enables SW-C level triggering of the same mechanisms as triggered by the Protection Hook.

The API looks as follows:

```
StatusType TerminateApplication
```

⁵ It is assumed that the integrator knows the mapping between Os-Application Ids (returned by the `GetApplicationId()` API), partitions, and the name of partitions (used to inform the RTE of a termination or request a restart).

```
(  
    ApplicationType Application,  
    RestartType      RestartOption  
)
```

Where `Application` refers to the ID of the OS-Application (generated during OS configuration), and `RestartOption` is either `RESTART` which will trigger the `OSRestartTask` (equivalent to `PRO_TERMINATEAPPL_RESTART` from the Protection Hook) or `NO_RESTART` which simply terminates the OS-Application (equivalent to `PRO_TERMINATEAPPL` from the Protection Hook).

10.3.1.4 OSRestartTask

The `OSRestartTask` is a task which is started when the Protection Hook returns `PRO_TERMINATEAPPL_RESTART`. When this task is executing it will be the only task running in the partition. The contents of the `OSRestartTask` are not standardized by AUTOSAR.

In the `OSRestartTask`, one can perform any cleanup activities necessary for being able to terminate or restart a partition. It should be noted that even though the task is named `OSRestartTask`, there is no automatic restart of the OS-Application (or the OS for that matter). During the execution of the `OSRestartTask`, there is still an active decision to make whether to terminate or restart the OS-Application (and thereby the partition).

10.3.1.5 OS-Application state machine

An OS-Application can be in one of the following states (see OS SWS [3] and Figure 2):

- **APPLICATION_ACCESSIBLE:** In this state, no protection violations have been detected and the software in the OS-Application runs normally.
- **APPLICATION_RESTARTING:** The OS-Application is in a state of emergency after the detection of a serious error. No resources of the OS-Application can be accessed from other OS-Applications and only the `OSRestartTask` is running. All other tasks/OsSrs/Alarms/etc. are terminated. The only way to get information into the OS-Application is to poll from the inside (i.e. from inside the `OSRestartTask`).
- **APPLICATION_TERMINATED:** The OS-Application is no more. Nothing is running. No access allowed from other OS-Applications. The only way to exit this state is to restart the OS (by restarting the ECU).

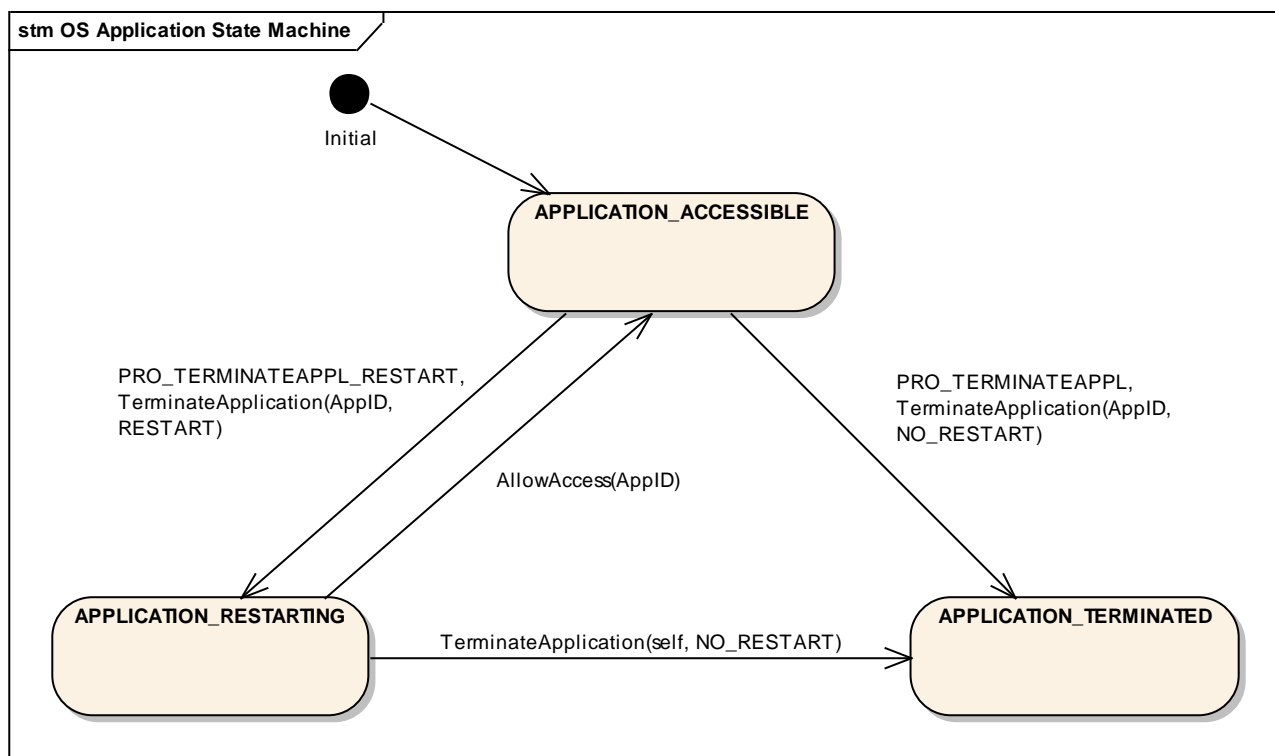


Figure 2: OS-Application State Machine

The OS-Application is set in the **APPLICATION_ACCESSIBLE** state by the `StartOS()` call. Once it has entered this state, all OS internal setup activities are completed, and the OS-Application is ready for execution. All state changes after startup are due to exceptional events which trigger the Protection Hook, or explicit requests to change the state of an OS-Application by calls to `TerminateApplication()`.

After the Protection Hook returns `PRO_TERMINATEAPPL_RESTART` to the OS, or the application-level request for restart is issued, the state of the OS-Application changes from **APPLICATION_ACCESSIBLE** to **APPLICATION_RESTARTING**. This means that resources held by the OS-Application are made inaccessible to other OS-Applications and no tasks belonging to the OS-Applications are running (or being activated). The OS then starts the `OSRestartTask` which can perform any necessary cleaning (e.g., in the SW-Cs and the BSW) before a decision is made to either restart or completely terminate the OS-Application.

Going from **APPLICATION_RESTARTING** to **APPLICATION_ACCESSIBLE** is achieved by allowing access to OS-Application resources using the call `AllowAccess()` in the `OSRestartTask`.

Going from **APPLICATION_RESTARTING** to **APPLICATION_TERMINATED** is achieved by an explicit request for termination from within the `OSRestartTask`, i.e., a call to `TerminateApplication()` with the ego-ID and `NO_RESTART` as parameters.

Going from **APPLICATION_ACCESSIBLE** to **APPLICATION_TERMINATED** is achieved by returning `PRO_TERMINATEAPPL` from the Protection Hook or by an external request using the `TerminateApplication()` call with `NO_RESTART` as

parameter. This transition should not be used when dealing with partitions as no cleanup activities will be performed here and thus the partition is potentially left in an inconsistent state which may persist in future starts of the ECU.

When in `APPLICATION_RESTARTING` calling `TerminateApplication()` with `ego-ID` and `RESTART` does not make sense. In fact may lead to an infinite loop (a new `OSRestartTask` is started, which may again call `TerminateApplication` with the `ego-ID` and `RESTART`).

10.3.2 Going from OS-Applications to partitions

At the OS level, an OS-Application can be terminated or restarted, as described above. Taking the step to partitions, we need to recognize that a partition is based on OS-Applications, but also encompasses BSW resources as well as SW-Cs, and may potentially have links and dependencies to external resources such as CDDs. Simply terminating or restarting an OS-Application will not be sufficient for terminating or restarting a partition: cleanup activities in the BSW as well as the SW-Cs (and potentially CDDs) may be necessary.

The description provided in this section uses the Protection Hook, the `OSRestartTask` and suitable functionalities in the RTE and BSW to implement the termination and restart of partitions. It should be noted that this is only one example of how this functionality can be incorporated – other solutions may be defined. The error handling strategy is not standardized by AUTOSAR. Using the basic mechanisms provided by AUTOSAR, different approaches to error handling could be implemented.

This description presents a suggestion for a partition state machine, usage of both Protection Hook and `OSRestartTask` and relevant calls to RTE and BSW, as well as necessary glue to get termination/restart working with AUTOSAR.

Error handling at partition level starts with the detection of an error which is serious enough to trigger a termination or restart of the said partition. Two sources for such errors are possible in AUTOSAR: 1) The OS detects (using HW support) protection violations (either memory access violation or timing protection violation) performed by a partition, corresponding to UC1 in Section 10.2.1, and 2) A SW-C can also detect (or be notified) of an error which requires a partition to be terminated or restarted and then requests this action by the OS, corresponding to UC2 in Section 10.2.2. A dedicated API exists for this and can also be used to terminate the own Partition (suicide).

10.3.2.1 Partition state machine

During its lifetime, a partition can be in a number of different states with respect to error handling⁶. These states are described in Figure 3.

⁶ The state machine described here is not provided by the AUTOSAR BSW. It is used here as a vessel for explaining the various states of a partition and which activities and conditions apply in these states

Figure 3 shows the states and transitions, and Table 10-1 and Table 10-2 give the corresponding state information and actions, and triggering conditions, respectively.

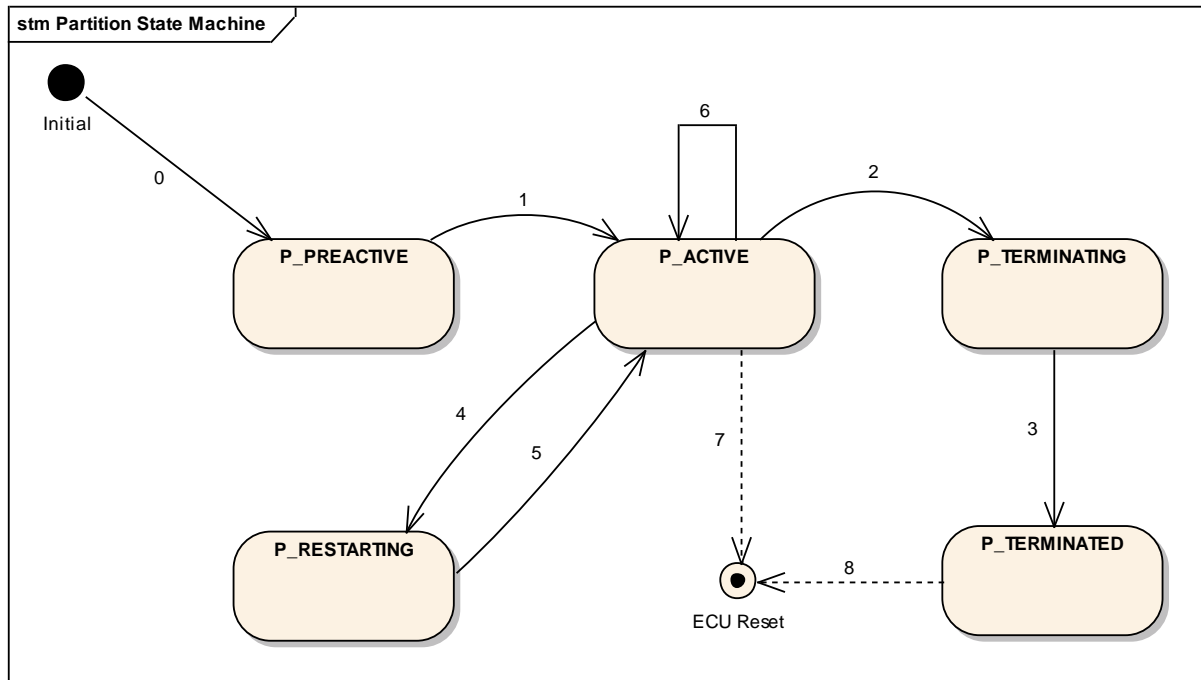


Figure 3: Partition state machine

The state machine has two stable states, **P_ACTIVE** and **P_TERMINATED**. Additionally, the transient states **P_TERMINATING** and **P_RESTARTING** are where the termination and restarting activities are being carried out. The **P_PREACTIVE** state indicates that the partition is not yet operational. This should be interpreted as a form of *UNDEFINED* state. Note that the state machine only captures the error handling state at system level of the partition; all other states and modes related to the partition and the SW-Cs within it are encapsulated in the **P_ACTIVE** state. This state machine is linked to the overall ECU state machine, since an ECU reset triggers a “restart” of the state machine.

In the state machine, we have included an entity called “ECU reset” to illustrate that a controlled ECU Reset (i.e. managed by EcuM) is possible only in the stable states **P_ACTIVE** and **P_TERMINATED**. It should be noted however, that an uncontrolled ECU reset (triggered via the reset pin on the micro controller) can occur in any state.

Table 10-1: The states of a partition w.r.t. Error Handling.

Partition state	OS-Application state	Description
P_PREACTIVE	N/A	Partition is not yet up and running.

and the transitions between them. If this state machine is to be implemented, this should be designed specifically for a given ECU.

		<p>All activities performed during normal ECU Startup are being performed. The RTE part of each partition is started, BSW resources are initialised, OS-Applications are set up, etc. These activities are the normal start-up activities of the ECU.</p>
P_ACTIVE	APPLICATION_ACCESSIBLE	<p>Partition running normally.</p> <p>In this state, the partition adheres to the overall ECU state. The partition remains in the P_ACTIVE state also while the Protection Hook is executed as a result of a protection violation.</p> <p>Upon entry into this state, the following needs to be done:</p> <ul style="list-style-type: none"> • RTE performs actions necessary to start partition execution (init runnables, activating schedule tables, etc.).
P_TERMINATING	APPLICATION_RESTARTING	<p>Partition is being terminated.</p> <p>In this state, some of the following actions may have to be performed:</p> <ul style="list-style-type: none"> • Report to DEM. Note that this is suggested to be done by the Protection Hook or the caller of TerminateApplication(), not by the OSRestartTask. • OS-Application is in the state APPLICATION_RESTARTING . Access from other OS-Applications to resources in the terminating OS-Application is not allowed by the OS. • Avoid shutdown, sleep, or reset of the ECU by requesting a certain mode to EcuM-Fixed or to the BswM. While a partition is in this state, the

		<p>ECU must not be shut down.</p> <ul style="list-style-type: none"> • Notify COM to stop communication for relevant signals. • Issue “No Communication” to ComM on behalf of the partition. • Stop relevant supervised entities in WdgM (according to configuration). • Invalidate relevant NvM-blocks (according to configuration). • Cancel pending jobs in the NvM on relevant NvM blocks (according to configuration). • Notify relevant CDDs that the partition is terminated such that they can perform necessary clean-up. • Notify IOHW abstraction layer that the partition is terminated, such that appropriate actions can be taken. • RTE shall ensure communication consistency (Section 10.3.5). • The RTE shall not activate any tasks in a terminating partition. • Trigger transition of the partition to the P_TERMINATED state. • Trigger transition of the OS-Application to the APPLICATION_TERMINATED state.
P_TERMINATED	APPLICATION_TERMINATED	<p>Partition is terminated and will no longer execute any SW-Cs until the ECU is reset.</p> <p>RTE ensures communication consistency (Section 10.3.5).</p>
P_RESTARTING	APPLICATION_RESTARTING	<p>The partition is being restarted and will return to the P_ACTIVE state when finished.</p>

		<p>In this state, some of the following actions may have to be performed:</p> <ul style="list-style-type: none"> • Report to DEM. Note that this is suggested to be done by the Protection Hook or the caller of TerminateApplication(), not by the OSRestartTask. • Partition is restarted at OS-level → The end results shall be an OS-Application (partition) ready-to-be-started by the RTE. Access from other OS-Applications to resources in the restarting OS-Application is disallowed. • Notify RTE to perform partition-specific clean-up and re-initialization. • Avoid shutdown, sleep, or reset of the ECU by requesting a certain mode to EcuM-Fixed or to the BswM. While a partition is in this state, the ECU must not be shut down. • Notify COM to restart communication for relevant signals. • Issue “No Communication” to ComM on behalf of the partition • Reset relevant supervised entities in WdgM. • Cancel pending jobs in the NvM on relevant NvM blocks (according to configuration). • Notify relevant CDDs that the partition is restarting such that they can perform necessary re-initialization. • Notify IOHW abstraction layer that the partition is restarting, such that appropriate actions can be taken. • RTE shall ensure communication consistency (Section 10.3.5). [RTE] • RTE shall not activate any
--	--	---

		<p>tasks in a terminated partition.</p> <ul style="list-style-type: none"> • Trigger transition of the partition to the P_ACTIVE state. [BswM] • Trigger transition of the OS-Application to the APPLICATION_ACCESSIBLE state which will notify the OS to once again allow access from other OS-Applications. <p>NOTE: If an application shall be restartable, the application has to take care of its NvM blocks itself. That is, it must perform consistency checks and any refresh or write-back activities necessary after a restart.</p>
ECU Startup	N/A	Normal ECU startup. Not affected by error handling. All partitions are set in the P_PREACTIVE state.
ECU Reset	N/A	ECU reset (here this also includes a complete shut-off and cold start of the ECU) will set the partition in the P_ACTIVE state after a normal startup. This is the only way to leave the P_TERMINATED state. Note that this means that even after a sleep/wakeup, the partition remains in the P_TERMINATED state.

Table 10-2: The transitions between states in the error handling state machine.

Transition	Description
0 ECU Startup → P_PREACTIVE	When the ECU begins its startup, all partitions are in the P_PREACTIVE state indicating that they are not yet up and running.
1 P_PREACTIVE → P_ACTIVE	When RTE is started, the partition can move to the P_ACTIVE state. Triggered when OS and RTE startup is finished.

<p>2</p> <p>P_ACTIVE → P_TERMINATING</p>	<p>An event has occurred that initiates the termination of a partition:</p> <ul style="list-style-type: none"> • The OS receives <code>PRO_TERMINATEAPPL_RESTART</code> from a call to the Protection Hook. • A call to <code>TerminateApplication()</code> is made with parameter set to <code>RESTART</code>. <p>The OS triggers the <code>OSRestartTask</code>. In that task, all actions necessary for terminating the partition are initiated.</p> <p>Note: The OS cannot distinguish between this transition and the transition <code>P_ACTIVE → P_RESTARTING</code>. This distinction needs to be communicated to the <code>OSRestartTask</code> by other means.</p>
<p>3</p> <p>P_TERMINATING → P_TERMINATED</p>	<p>When all clean-up activities related to termination are completed the partition is considered terminated. The transition to the <code>P_TERMINATED</code> state is linked with the transition of the OS-Application from <code>INACCESSIBLE</code> to <code>TERMINATED</code>.</p>
<p>4</p> <p>P_ACTIVE → P_RESTARTING</p>	<p>An event has occurred that initiates the restart of the partition:</p> <ul style="list-style-type: none"> • The OS receives <code>PRO_TERMINATEAPPL_RESTART</code> from a call to the Protection Hook. • A call to <code>TerminateApplication()</code> is made with the parameter set to <code>RESTART</code>. <p>The OS triggers the <code>OSRestartTask</code>. In that task, all actions necessary for terminating the partition are initiated.</p> <p>Note: The OS cannot distinguish between this transition and the transition <code>P_ACTIVE → P_TERMINATING</code>. This distinction needs to be communicated to the <code>OSRestartTask</code> by other means.</p>
<p>5</p> <p>P_RESTARTING → P_ACTIVE</p>	<p>When all clean-up activities related to termination are completed the partition can go back to the <code>P_ACTIVE</code> state.</p> <p>The transition to the <code>P_ACTIVE</code> state is linked with the transition of the OS-Application from <code>APPLICATION_RESTARTING</code> to <code>APPLICATION_ACCESSIBLE</code>.</p>
<p>6</p> <p>P_ACTIVE → P_ACTIVE</p>	<p>Protection Hook is executed and returns <code>PRO_IGNORE</code>.</p>

7 P_ACTIVE ECU Reset	→	An event occurred that initiates the shutdown of the OS (i.e., ECU), i.e., the OS receives a <code>PRO_SHUTDOWN</code> from the Protection Hook. This will eventually lead to ECU reset. This could also be the result of an ECU shutdown triggered by another source (external watchdog) and managed by the EcuM.
8 P_TERMINATED ECU Reset	→	After an ECU reset the partition will move to the P_ACTIVE state after normal start-up. This is the only transition for the partition from the P_TERMINATED state. Note that this means that, even after sleep/wakeup a terminated partition is still terminated.

10.3.2.2 Error handling strategy in the Protection Hook

The Protection Hook implements the overall error handling strategy for the partition related to protection violations. When the OS detects a protection violation, the Protection Hook will be called and a decision on proper reaction to the detected error can be made.

The following information should be gathered to form the basis for the decision on how to react to the violation:

- **Partition ID:** in which partition has the protection violation occurred? There is only one Protection Hook for the entire ECU, so the partition ID is necessary to implement partition specific error handling strategies. The partition ID can be found using `GetApplicationID()`. This call will return the identifier of the OS-Application to which the offending task belongs.
- **Error type:** what was the reason for the protection violation? Depending on the type of violation, different reactions may be selected. The error type is provided as a parameter to the Protection Hook.
- **Partition restart counter:** how many times has the partition been restarted (since the last ECU startup)? If a partition has been restarted a certain number of times (pre-defined threshold) it may be more suitable to terminate the partition instead of attempting another restart. The restart counter needs to be implemented by the integrator/developer, e.g., using a static structure to store the counters of the relevant partitions.
- **OS-Application State:** what is the state of the offending OS-Application? If the OS-Application is in the `APPLICATION_RESTARTING` state, this is an indication that there is an error in the error handling. This can only occur if a protection violation occurs while executing the `OSRestartTask`.

This information allows the integrator/developer to define an error handling strategy which takes into account relevant information about the partition, the error and the error handling history.

Based on the input the Protection Hook needs to decide what to do:

- **Do nothing.** If the cause of the protection violation was identified as transient and benign and continued execution is not foreseen to experience any

problems, this may be a viable option. For this, the Protection Hook shall return `PRO_IGNORE`.

- **Restart the partition.** If the cause of the protection violation was identified transient and a re-initialisation of the partition is foreseen to ensure that execution will not experience any problems, this may be a viable option. For this, Protection Hook shall return `PRO_TERMINATEAPPL_RESTART` as this will trigger the `OSRestartTask` where additional clean-up/re-initialisation of the BSW and SW-Cs will be performed. Note that the same return code is used for termination of a partition, so the decision has to be communicated to the `OSRestartTask` in some other way.
- **Terminate the partition.** If the cause of the protection violation is deemed permanent, or a re-initialisation is not seen as being a successful reaction, termination may be a viable option. For this, Protection Hook shall return `PRO_TERMINATEAPPL_RESTART` as this will trigger the `OSRestartTask` where additional clean-up/re-initialisation of the BSW and SW-Cs will be performed. Note that the same return code is used for termination of a partition, so the decision has to be communicated to the `OSRestartTask` in some other way.
- **Shutdown OS/ECU.** If the cause of the protection violation is deemed so severe that the only alternative is to shut down the ECU, then this is done by return `PRO_SHUTDOWN`. This will trigger a shutdown of the OS which in turn will eventually shut down the entire ECU. The `OSRestartTask` is not triggered.

The Protection Hook may also return `PRO_TERMINATEAPPL`. This will terminate the OS-Application at the OS level but will not trigger the `OSRestartTask`. Therefore, there is no way to perform clean-up activities in the BSW or SW-Cs or elsewhere and the risk for an inconsistent state in the ECU is very high. Thus, in the approach described here, this return code is not used.

The option to terminate only the offending task by returning `PRO_TERMINATETASK` shall not be used, since it does not terminate the whole partition. This way, consistency issues only arise in inter-partition communication, where they are more easily identified and handled.

When a decision has been made on the proper reaction to the error, a report must be made to the DEM to log this event. If this report is not made by the Protection Hook, the decision needs to be communicated to the reporting entity. Note that the `OSRestartTask` will only be triggered if the return value is `PRO_TERMINATEAPPL_RESTART`.

Before returning to the OS and thereby notifying the OS of the decision the Protection Hook must notify the RTE such that the RTE can guarantee data and communication consistency across partitions (and ECUs). This is done with the call `Rte_PartitionTerminated_<partition>()` if the partition is to be terminated or `Rte_PartitionRestarting_<partition>()` if the partition is to be restarted.

10.3.2.3 Clean-up activities in the `OSRestartTask`

In the `OSRestartTask` the integrator/developer places all the clean-up and notification activities necessary for terminating or restarting a partition.

As the OSRestartTask is triggered for restarting as well as terminating partitions, the first thing to do is to find out what the required actions shall be. There are no standard AUTOSAR mechanisms to communicate the decision from the Protection Hook to the OSRestartTask, so this would have to be implemented explicitly. A suggested solution is to implement the presented partition state machine and let the OSRestartTask check the state variable holding the current state of the partition. If it is P_TERMINATING the partition shall be terminated and if it is P_RESTARTING it shall be restarted.

Once it is known whether the partition shall be restarted or terminated, proper clean-up actions can be performed. Table 10-3 contains some actions which may have to be performed for restart and termination.

Module	Clean-up activities
COM	<p>Terminate:</p> <ul style="list-style-type: none"> Issue a request to stop sending signals (I-PDU groups) for the terminated partition. <p>Restart:</p> <ul style="list-style-type: none"> Issue a request to stop sending signals (I-PDU groups) for the partition being restarted. After other cleanup activities are finished, sending of signals is started again and their values are re-initialized.
ComM	<p>Terminate:</p> <ul style="list-style-type: none"> Issue a “No communication” request on behalf of the partition to ComM. <p>Restart:</p> <ul style="list-style-type: none"> Issue a “No communication” request on behalf of the partition to ComM.
DCM	<p>Terminate:</p> <ul style="list-style-type: none"> No explicit clean-up necessary. However, calls from DCM to terminated partitions will be left unanswered. The RTE will indicate a time-out to the caller. <p>Restart:</p> <ul style="list-style-type: none"> No explicit clean-up necessary. However, calls from DCM to will be left unanswered during restart. The RTE will indicate a time-out to the caller.

Module	Clean-up activities
DEM	<p>Terminate:</p> <ul style="list-style-type: none"> No explicit clean-up necessary. However, monitors contained in terminated partitions need to be managed, e.g., shut off. <p>Restart:</p> <ul style="list-style-type: none"> Run-time information of monitors located in the restarted partition needs to be managed.
EcuM	<p>Terminate:</p> <ul style="list-style-type: none"> Avoid shutdown, sleep, or reset of the ECU by requesting a certain mode to EcuM-Fixed or to the BswM. When clean-up is completed, this request can be cancelled. <p>Restart:</p> <ul style="list-style-type: none"> Avoid shutdown, sleep, or reset of the ECU by requesting a certain mode to EcuM-Fixed or to the BswM. When clean-up is completed, this request can be cancelled. The SW-Cs of the partition will issue similar requests when execution commences.
NvM	<p>Terminate:</p> <ul style="list-style-type: none"> Invalidate relevant blocks belonging to the partition (note that not all blocks of a partition may require invalidation). Outstanding requests from the terminated partition shall be removed from the job queue. <p>Restart:</p> <ul style="list-style-type: none"> No impact. Restart activities and consistency checks for NvM must be handled by the SW-Cs. This may include invalidating invalid blocks after a restart. <p>Note: An integrator may decide which blocks can be “trusted” also in case of a termination. This makes different policies for “normal” shutdown and “exceptional” shutdown of partitions possible.</p>

Module	Clean-up activities
WdgM	<p>Terminate:</p> <ul style="list-style-type: none"> Stop all supervised entities belonging to the partition. <p>Restart:</p> <ul style="list-style-type: none"> The integrator has to choose either to temporarily stop the supervision during restart or configure the watchdog by taking the time to restart the partition into account
RTE	<p>Terminate:</p> <ul style="list-style-type: none"> The local RTE of the partition has already been notified by the Protection Hook that the partition is being terminated. No further actions necessary. <p>Restart:</p> <ul style="list-style-type: none"> Once all clean-up activities in the BSW have been completed and the partition is ready for execution again, the RTE needs to be restarted.
IoHwAb	<p>Terminate:</p> <ul style="list-style-type: none"> Set relevant IOs of IoHwAb to proper values. This is likely to be very application specific. <p>Restart:</p> <ul style="list-style-type: none"> Reset relevant IOs of IoHwAb to proper values. This is likely to be very application specific.
CDD	<p>Terminate:</p> <ul style="list-style-type: none"> Notify relevant CDDs that the partition is being terminated. The CDDs will then have to take appropriate actions. <p>Restart:</p> <ul style="list-style-type: none"> Notify relevant CDDs that the partition is being restarted. The CDDs will then have to take appropriate actions

Table 10-3. Clean-up activities for BSW

All clean-up activities in BSW must be initiated with `CallTrustedFunction()` in order to switch to the BSW OS-Application, since the `OSRestartTask` is running in an OS-Application that is currently in the state `APPLICATION_RESTARTING`.

Note that clean-up activities in BSW modules are not necessarily synchronous. Thus, there may be a need to wait for all activities to be completed.

Once all relevant BSW clean-up activities are completed, the OSRestartTask can perform the final activities:

- If the partition is to be restarted
 - Request a restart of the partition local RTE by calling `Rte_RestartPartition_<partition>()`.
 - Request the OS to allow access to the OS-Application by calling `AllowAccess()`. This will set the OS-Application state to `APPLICATION_ACCESSIBLE`.
 - Set the partition state to `P_ACTIVE`.
 - Call `TerminateTask()` to terminate the OSRestartTask.
- If the partition is to be terminated
 - Set the partition state to `P_TERMINATED`.
 - Call `TerminateApplication(self, NO_RESTART)` which will terminate the OS-Application without triggering the OSRestartTask again. This will set the OS-Application in the `APPLICATION_TERMINATED` state.

10.3.2.4 Externally triggered restart or termination

In UC2, a trusted SW-C can trigger a termination/restart of another partition. Such an external trigger requires the following actions to be made in the SW-C:

1. Notify RTE whether the partition will be terminated using the call `Rte_PartitionTerminated_<partition>()` or will be restarted using the call `Rte_PartitionRestarting_<partition>()`.
2. Make sure to post the decision (termination or restart) such that the OSRestartTask can decide which actions to take once the clean-up activities commence.
3. Initiate termination/restart of the partition with the call `TerminateApplication(<partition>, RESTART)`. The value `RESTART` is needed to trigger the OSRestartTask so that clean-up activities can be performed.
4. Report to DEM.

From this point, the actions will be the same as when the Protection Hook returns its decision.

Automotive applications can be distributed over several partitions and thus over several ECUs. The `TerminateApplication()` call only covers local partitions (i.e. on the same ECU as the caller), so the handling of distributed applications needs to be considered at application level.

One way to implement a distributed monitoring and control of related partitions is to have a dedicated application-level error manager (ALEM) in each ECU. The ALEM will monitor and control the partitions in an ECU and can communicate relevant information and requests for termination or restart of remote partitions to the ALEMs concerned.

10.3.3 Sequence diagram for termination and restart of a partition

The sequence of events for terminating or restarting a partition is shown in Figure 4. The details are described in Section 10.3.2.

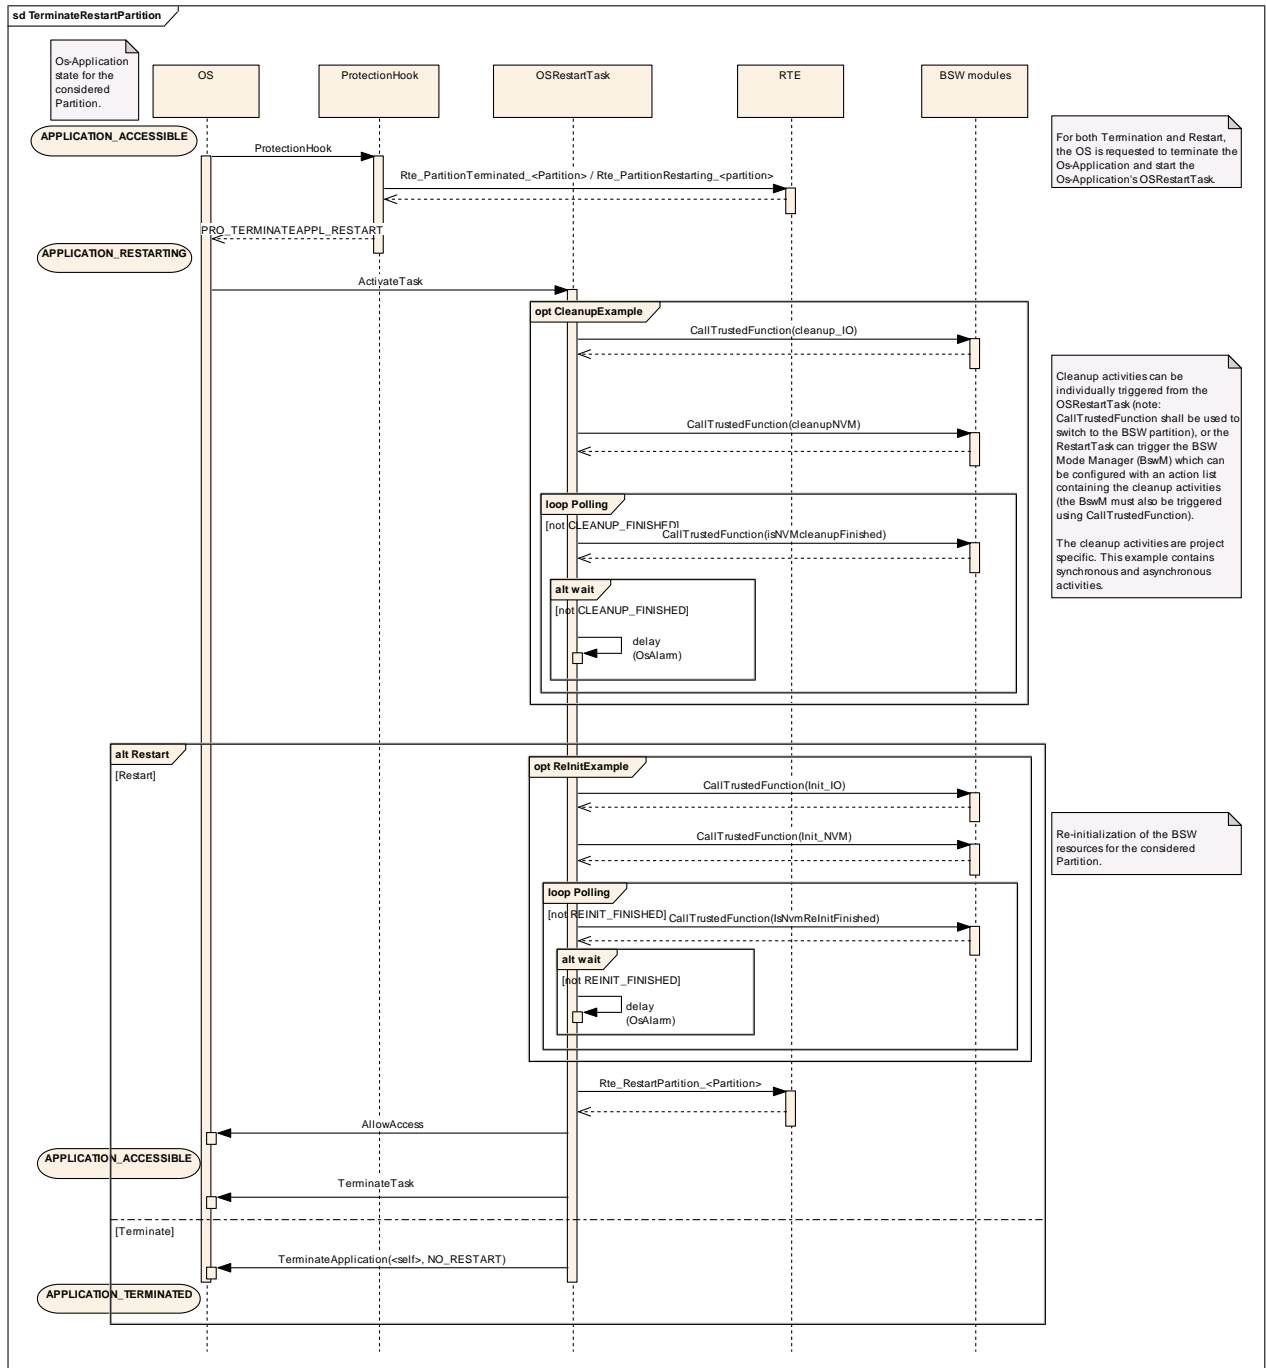


Figure 4: Sequence diagram showing the interactions when terminating or restarting a partition.

10.3.4 Support for Use Cases

10.3.4.1 UC1: SW Partitioning

By placing each application requiring protection in an OS-Application, partitioning between applications is achieved. Possible reactions to errors include terminating and restarting of the affected partition. UC1 is therefore fully covered.

Applications may also be spread across multiple partitions to provide more fine-grained protection, but this may incur communication overhead during run-time, due to communication across partition boundaries.

10.3.4.2 UC2: OEM Specific Error Handling

An OEM requiring specific error handling policies needs to place applications in different partitions. Furthermore, if forcible termination and restart of applications (i.e., partitions) is required, dedicated trusted SW-Cs can request this by the OS. UC2 is therefore fully covered.

10.3.5 Consistency Aspects

10.3.5.1 Application-level consistency

Any application-level inconsistencies arising from restarting a partition are expected to be handled by the SW-Cs themselves. However, SW-Cs can trust the RTE to behave consistently and in a defined way w.r.t. terminated/restarting partitions. Section 10.3.5.3 presents the consistency requirements needed for the RTE to achieve this.

10.3.5.2 BSW consistency

BSW consistency needs to be ensured by clean-up activities performed when restarting or terminating a partition. Details about these activities are described in Section 10.3.2.3.

10.3.5.3 Communication consistency

When a partition is terminated or restarted, this may lead to a (temporary) loss of communication between communicating parties. In all such cases, AUTOSAR must ensure that the system is consistent, i.e., the system can still make progress without undesired side-effects. It is the job of the application developers to handle cases where SW-Cs (contained in the partitions) are terminated or restarted.

To handle the case where a partition is terminated and consequently no longer able to react to events, the behaviour of the system w.r.t. the terminated partition must be defined. Note that since the error containment region is the partition, consistency for intra-partition communication is irrelevant, since all tasks in the partition are terminated. Thus only inter-partition (superset of inter-ECU) communication can lead to consistency problems.

The main principle shall be as follows: Consider the case where we have two partitions interacting with each other. At some point, one of the partitions is terminated (restarted). For the remaining partition, the perception of this shall be the same regardless of whether the terminated (restarted) partition resides on the same ECU or another ECU. As a consequence, the RTE behaviour shall be the same in both cases.

This principle implies that timeout monitoring shall be used to detect if the receiver of information (both in SR and CS communication) is terminated. Trying to initiate communication with a partition which is terminated (or being restarted) shall result in a timeout notification, the same as if the original signal had been lost or the receiving ECU did not respond. Also, it does not impose on the partition to “know” if communication is local or remote, even in case of failure. After a partition is restarted its state might be inconsistent with other partitions’ view of the state. For instance, if there is a dependency between two messages (both CS & SR) it may be broken by a restart.

10.4 Integrator Responsibility

The integrator has the overall responsibility for implementing the terminate and restart actions required by providing the code for the Protection Hook and the Restart Task as well as any glue code necessary for coordination and proper incorporation of the functionality.

The following is a checklist of things that need to be done by the integrator or developer in order to get proper handling of termination and restart of partitions:

1. Decompose applications into SW-Cs and group the SW-Cs according to chosen partitions/error containment regions.
2. Configure partitions and OS-Applications in accordance with outcome of decomposition.
3. Indicate for each SW-C whether it can be terminated or restarted (and of course, make sure the implementation of the SW-C supports this).
4. Indicate for each partition whether it can be terminated or restarted.
5. Decide upon error handling policies and strategies for the partitions. Base decision of restart and termination of a given partition on e.g. error type and number of previous restarts, and any other information relevant for the decision.
6. Implement the chosen policies and strategies in the Protection Hook (remember that there is one global Protection Hook for the entire ECU). Notify the RTE of the decision using `Rte_PartitionTerminated_<partition>()` or `Rte_PartitionRestarting_<partition>()`. Return the proper code representing the decision.
 - `PRO_IGNORE` if nothing shall be done.

- `PRO_TERMINATEAPPL_RESTART` if the partition shall be restarted or terminated.
 - `PRO_SHUTDOWN` if the OS (and thus eventually the ECU) shall be shut down.
7. Implement the necessary clean-up actions in the `OSRestartTask` of the partition (there is one such task per partition) and configure the BSW accordingly. Remember that this task should be triggered both when terminating and when restarting a partition, and that the necessary clean-up actions are not necessarily the same for both activities. Thus, the `OSRestartTask` must be able to find out whether the partition shall be terminated or restarted in order to perform the correct clean-up actions. The suggested way to do this is to implement the partition state machine described in Section 10.3.2.
 8. If there is a need for application-level coordination of error handling across multiple partitions, implement an application-level error manager (ALEM) either as a dedicated SW-C or as part of another SW-C. Note that such an ALEM must not be in any of the partitions it is set to monitor. Use the `TerminateApplication()` API to control the termination and restart of partitions from an ALEM.